

# Canonicalizing Execution for Automatic Debugging

Xiangyu Zhang



December 16, 2011 @ ISHCS, PKU

# Students

---



- ❑ William Nick Sumner (expected May 2012)  
and Bin Xin (Google)





# Problem Statement

---

- ❑ Given a software runtime failure, we aim to automatically compute the explanation of the failure
  - A chain of execution steps that are causally correlated, leading from the root cause to the failure, called **the causal path of the failure**



# A Conceptual Example

```
inventory = [(Shoes,5); (Hats,0); (Ties,1)]  
bought = 0
```

```
for (item, available) in inventory:  
    if bought < 3 and available >= 0:  
        buy(item)  
        bought += 1
```

```
print "Items bought: ", bought
```

Should print "Items bought: 2"

**Failure:** prints "Items bought: 3"



# The Causal Path

A faulty branch is taken at **A** ;  
so bt is given the faulty value 2 at **B** ;  
so bt is given the faulty value 3 at **C** ;  
so '3' is printed erroneously at **D** .

1. Leading from the root cause to the failure;
2. Values along the path are faulty;
3. Steps are causally correlated.

```
1) for (itm, av) = (S,5):  
2)  if bt < 3 and av >= 0:  
3)    buy(itm)  
4)    bt += 1  
1) for (itm, av) = (H,0):  
2)  if bt < 3 and av >= 0:  
3)    buy(itm)  
4)    bt += 1  
1) for (itm, av) = (T,1):  
2)  if bt < 3 and av >= 0:  
3)    buy(itm)  
4)    bt += 1  
5) print bt
```



# The Basic Idea

---

- ❑ Acquire a passing execution as the reference
- ❑ Compare the states of the passing and failing runs to isolate the failure causal path
- ❑ Delta Debugging (by Andreas Zeller) was proposed in 2002 to compute failure causal path
  - Our algorithm is completely different
  - We solved a number of key challenges for practicality

so '3' is printed erroneously at **D**.

Failing

```
1)for av = 5:
2) if ... and av >= 0:
3)  buy(item)
4)  bt += 1 //=1
1)for av = 0:
2) if ... and av >= 0:
3)  buy(item)
4)  bt += 1 //=2
1)for av = 1:
2) if ... and av >= 0:
3)  buy(item)
4)  bt += 1 //=3
D 5)print bt
```

Correct (patched)

```
1)for av = 5:
2) if ... and av >= 0:
3)  buy(item)
4)  bt += 1 //=1
1)for av = 0:
2) if !(... and av >= 0):
3)  buy(item)
4)  bt += 1 //=2
5)print bt
```

3 is printed vs. 2 is printed

so bt is given the faulty value 3 at **C** ;  
so '3' is printed erroneously at **D** .

Failing

```
1)for av = 5:
2) if ... and av >= 0:
3)  buy(item)
4)  bt += 1 //=1
1)for av = 0:
2) if ... and av >= 0:
3)  buy(item)
4)  bt += 1 //=2
1)for av = 1:
2) if ... and av >= 0:
3)  buy(item)
C 4)  bt += 1  //=3
D 5)print bt
```

Correct (patched)

```
1)for av = 5:
2) if ... and av >= 0:
3)  buy(item)
4)  bt += 1 //=1
1)for av = 0:
2) if !(... and av >= 0):
1)for av = 1:
2) if ... and av >= 0:
3)  buy(item)
4)  bt += 1  //=2
5)print bt
```

bt =3 vs. bt=2



so bt is given the faulty value 2 at **B** ;  
so bt is given the faulty value 3 at **C** ;  
so '3' is printed erroneously at **D** .

Failing

1)for av = 5:  
2) if ... and av >= 0:  
3) buy(item)  
4) bt += 1 //=1  
1)for av = 0:  
2) if ... and av >= 0:  
3) buy(item)  
**B** 4) bt += 1 //=2  
1)for av = 1:  
2) if ... and av >= 0:  
3) buy(item)  
**C** 4) bt += 1 //=3  
**D** 5)print bt

Correct (patched)

1)for av = 5:  
2) if ... and av >= 0:  
3) buy(item)  
4) bt += 1 //=1  
1)for av = 0:  
2) if **!(... and av >= 0):**  
1)for av = 1:  
2) if ... and av >= 0:  
3) buy(item)  
4) bt += 1 //=2  
5)print bt

bt =2 vs. bt=1

A faulty branch is taken at **A** ;  
so bt is given the faulty value 2 at **B** ;  
so bt is given the faulty value 3 at **C** ;  
so '3' is printed erroneously at **D** .

Failing

1)for av = 5:  
2) if ... and av >= 0:  
3) buy(item)  
4) bt += 1 //=1  
1)for av = 0:  
**A** 2) if ... and av >= 0:  
3) buy(item)  
**B** 4) bt += 1 //=2  
1)for av = 1:  
2) if ... and av >= 0:  
3) buy(item)  
**C** 4) bt += 1 //=3  
**D** 5)print bt



Correct (patched)

1)for av = 5:  
2) if ... and av >= 0:  
3) buy(item)  
4) bt += 1 //=1  
1)for av = 0:  
2) if **!(... and av >= 0)**:  
  
1)for av = 1:  
2) if ... and av >= 0:  
3) buy(item)  
4) bt += 1 //=2  
5)print bt

**(... and av >= 0) is True** vs. **!(... and av >= 0) is False**



# Algorithm Overview

- 1 align (E1, E2) *Execution alignment*
- 2 **FOR** each aligned point n in the reverse execution order **DO**
- 3     re-execute both E1 and E2 to n
- 4     take memory snapshots M1 and M2, of E1 and E2, resp.
- 5     canonicalize M1 and M2 *Memory canonicalization*
- 6      $\Delta$  = state differences between M1 and M2
- 7     report the minimal subset of  $\Delta$  critical to the failure
- 8 **ENDFOR** *Causality testing*

Test frame rejection is a by-product of the process flow by which  $\Delta$  is a sequence of differences from the failing run, which may include complex data structures and pointer values



# Outline

---

- ❑ Execution alignment
- ❑ Canonicalizing memory
- ❑ Causality testing
- ❑ Results



# Outline

---

- ❑ Execution alignment
- ❑ Canonicalizing values
- ❑ Causality testing
- ❑ Results



# Execution Alignment

---

- ❑ Informal problem statement
  - Given an execution point in a run, what is the corresponding point in another run
- ❑ Current practice
  - statement (+ instance)

Problematic in the presence of control flow differences, loops, and recursion.



# Instance Count Does NOT work

Failing

```
1)for ...  
2) if ... and av >= 0:  
4)  bt = 1
```

Correct (patched)

```
1)for ...  
2) if ... and av >= 0:  
4)  bt = 1
```

```
1)for ...  
2) if ... and av >= 0:  
4)  bt = 2
```

```
1)for ...  
2) if !(... and av >= 0):
```

```
1)for ...  
2) if ... and av >= 0:  
4)  bt = 3
```

```
1)for ...  
2) if ... and av >= 0:  
4)  bt = 2
```

```
5)print bt
```

```
5)print bt
```

bt =2 vs. bt=2

# Our Solution – Execution Indexing (PLDI'08)

---



## □ Concept

- Develop a canonical representation, called the **index**, of an execution point. Points with the same index across runs align

## □ Basic idea

- An execution is modeled as a trace. The trace is parsed into **an index tree** that represents its nesting structure. Indices are defined based on the tree.



1) **for** av = 5:

2) **if** ... **and** av >= 0:

3) buy(item)

4) bt = 1

1) **for** av = 0:

2) **if** ... **and** av >= 0:

3) buy(item)

4) bt = 2


1) **for** av = 1:

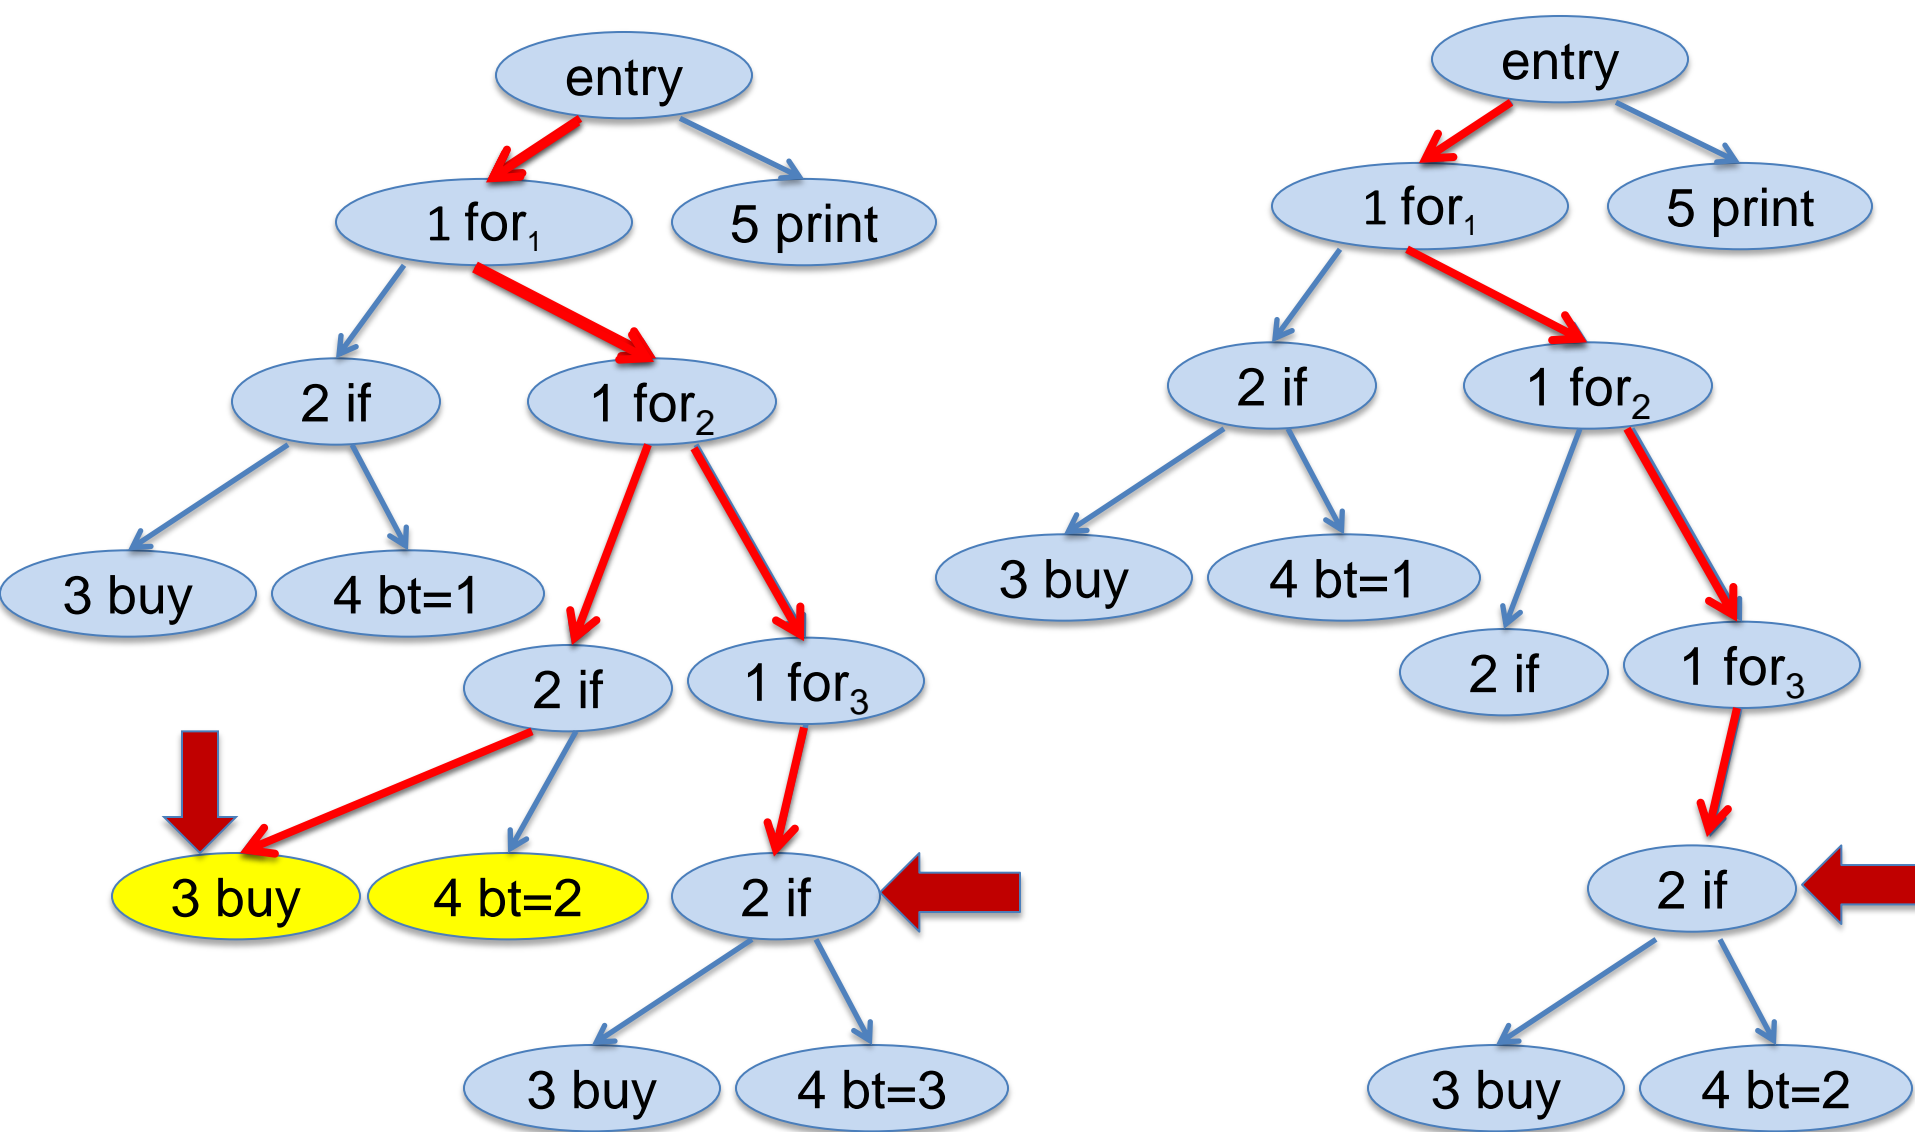
2) **if** ... **and** av >= 0:

3) buy(item)

4) bt = 3

5) **print** '3'





The index of an execution point is the path leading from the root to that point

Points with the same index align

**[entry, 1, 1, 2, 2]**



# Indexing

---

- ❑ We don't actually collect trace
- ❑ We compute the index of the current execution point on the fly
  - In order to make use of the online algorithm.
    - At a point of interest, issue the query “acquire the current index  $X$ ” (e.g.  $X=[\text{entry}, 1, 1, 1, 2]$ )
    - In the other run, issue “execute to the index  $X$ ”.



# The Indexing Algorithm

- ❑ Use a stack to represent the current index
- ❑ Push on a predicate, a function entry
- ❑ Pop on the immediate post-dominator or function exit.

[entry]	→	Entry
[entry, 1]	→	1)for ...
[entry, 1, 2]	→	2) if ... and av >= 0:
[entry, 1, 2]	→	4)    bt = 1
[entry, 1, 1]	→	1)for ...
		2) if ... and av >= 0:
		4)    bt = 2
		1)for ...
		2) if ... and av >= 0:
		5)print bt

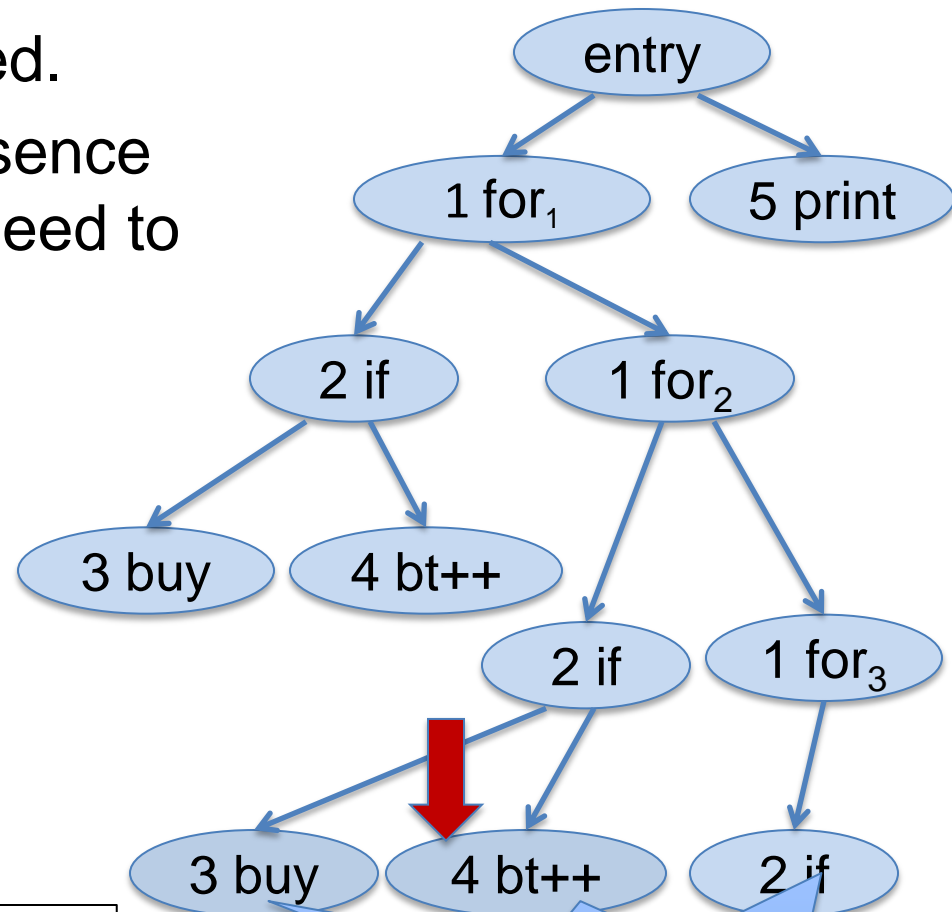


# Indexing Algorithm Continued

- ❑ Iterations are coalesced.
- ❑ Predicates whose presence can be inferred don't need to be pushed

[entry, 1, 1, 2, 4]  
↓  
[entry, 1:2, 2, 4]  
↓  
[entry, 1:2, 4]

```
1)for ...  
2) if ... and av >= 0:  
4)    bt += 1
```





# Outline


---

- ❑ Execution alignment [PLDI'08]
- ❑ **Canonicalizing memory [FSE'10]**
- ❑ Causality testing
- ❑ Getting the passing run
- ❑ Results




# Canonicalizing Memory

- Informal problem statement
  - The same heap object may be present in different locations in the two runs
  - Pointer values are not directly comparable



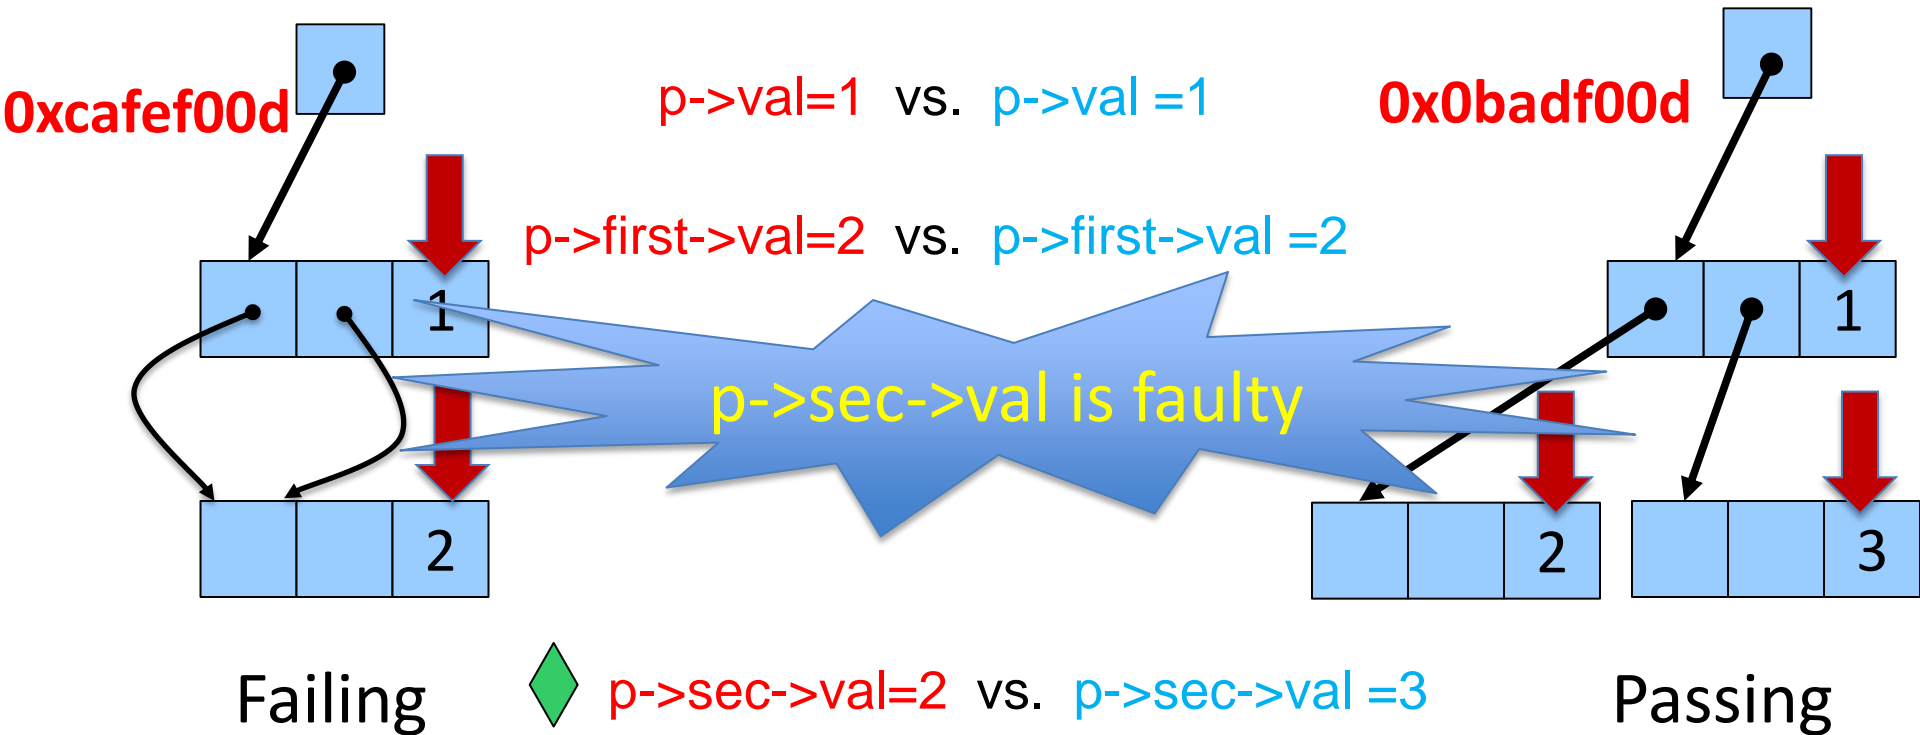
Failing



Passing



- ❑ Solution: heap object correspondence is established by symbolic reference path
  - imprecision due to aliasing





# Our solution – Memory Indexing



## □ Concept

- Develop a canonical representation, called the **memory index**, for each allocated heap location. Locations with the same memory index across runs align

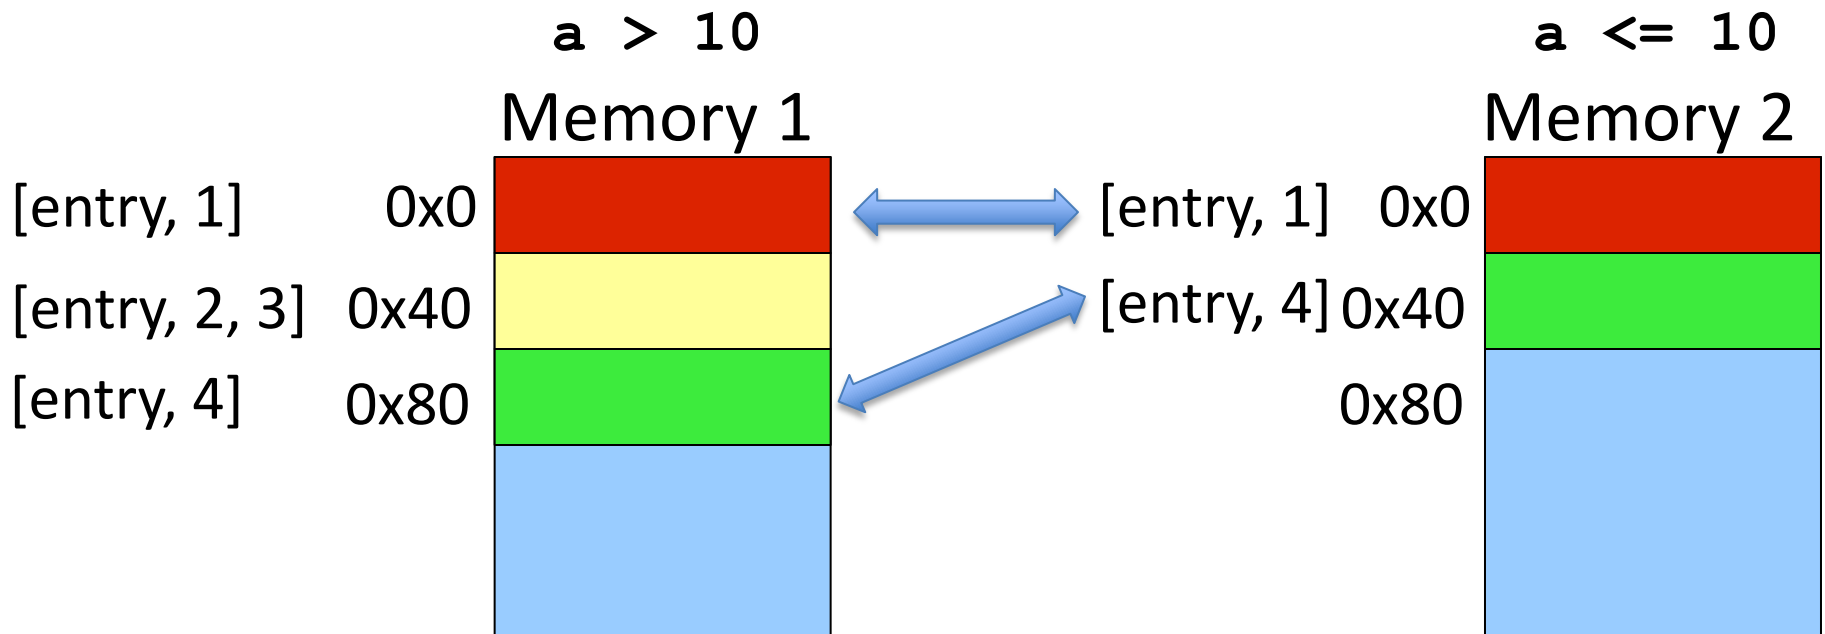
## □ Basic idea

- An allocated heap object is denoted by **the control flow index** of its dynamic allocation point
  - Memory alignment is reduced to control flow alignment
- A precise, dynamic analogue to static heap abstraction



# A Conceptual Example

```
1) x = malloc(64)
2) if (a > 10)
3)     y = malloc(64)
4) z = malloc(64)
```





# Online Algorithm – Flow Semantics

Execution

Action

`p=malloc (...)`

`mldx[p]=(current_cfidx,0)`

`q= p + c`

`mldx[q]= (mldx[p].first,  
mldx[p].second+ c)`

```
1  if (...)
2      if (...)
3          p= malloc(64)
4  q = p + 10
```

`mldx[p]=([entry,1,2,3], 0)`

`mldx[q]=([entry,1,2,3], 10)`



# Outline

---

- ❑ Execution alignment [PLDI'08]
- ❑ Canonicalizing memory [FSE'10]
- ❑ **Causality testing**
- ❑ Results



# Causality Testing

---

- ❑ Informal problem statement
  - At aligned points, determine whether a subset of state differences causes the failure, by testing whether copying the subset from the failing run to the passing run induces the failure
- ❑ Challenges
  - Copying state from one execution to another in a robust way
- ❑ Basic idea
  - Concrete memory  $\Rightarrow$  canonical memory
  - Compare canonical memory
  - Copy canonical differences from the failing to the pass
  - Concretize these differences in the passing run
  - Resume execution



# Outline

---

- ❑ Execution alignment [PLDI'08]
- ❑ Canonicalizing memory [FSE'10]
- ❑ Causality testing
- ❑ **Results**



# Implementation

---

- ❑ LLVM, Python, and C++
  - >10K LOC
- ❑ Support programs in C/C++
  - as large as gcc

# Effectiveness in Auto-Debugging



Program	Bug ID	Passing	Max Diff (w. MI)	Max Diff (w/o MI)	Step/ Diff	Time (s)
gcc-2.95	529	patching	233	8365	8/1	4559
gcc-2.95	776	patching	230	10101	2/1	379
gcc-2.95	2771	diff input	284	11095	4/1	1797
make 3.81	16958	non-reg	184	2699	9/1	740
make 3.81	18435	diff input	336	3301	29/2	645
make 3.81	19133	diff input	550	3728	5/2	235
make 3.80	112	diff input	645	3309	11/6	685
gawk 3.1	2006	diff input	22	630	8/1	56

- ❑ Canonicalization is very critical
- ❑ Computed causal paths are short and thin
- ❑ Computation time is reasonable
- ❑ Other over 20 real bugs for grep, gzip, tar, and flex.





# Recent Results

---

- ❑ Finishing validating our technique on a very popular debugging test-bed – Siemens suite
  - >10,000 runtime failures
  - Precisely capture the root causes
  - Present the precise explanations for all those failures online

# Demo

---



<http://www.youtube.com/watch?v=36cXwVqMo30>



# Related Work

---

- ❑ Delta debugging [Zeller, FSE '02]
- ❑ Statistical debugging [Liblit, PLDI'05]
- ❑ Dynamic slicing [Zhang, PLDI'04]
- ❑ Trace based correspondence [Ramanathan, FASE '06]
- ❑ Others...



# Conclusions

---

- ❑ We develop three important primitives
  - Control flow indexing
  - Memory indexing
  - Causality testing
- ❑ An effective auto-debugging engine can be composed from the primitives
  - Handles real bugs for non-trivial programs such as gcc
- ❑ Execution comparison needs canonicalization

Thank you!

<http://www.cs.purdue.edu/homes/xyzhang>