


# Modular Verification of Linearizability with Non-Fixed Linearization Points

Xinyu Feng  
USTC


Joint work with Hongjin Liang

# Concurrency Verification

- Concurrency: a basic programming skill in multicore era
- Very difficult to ensure correctness
  - Very subtle interleaving
- This talk:  
correctness of (non-blocking) concurrent obj.



# Example: Treiber's Non-blocking Stack



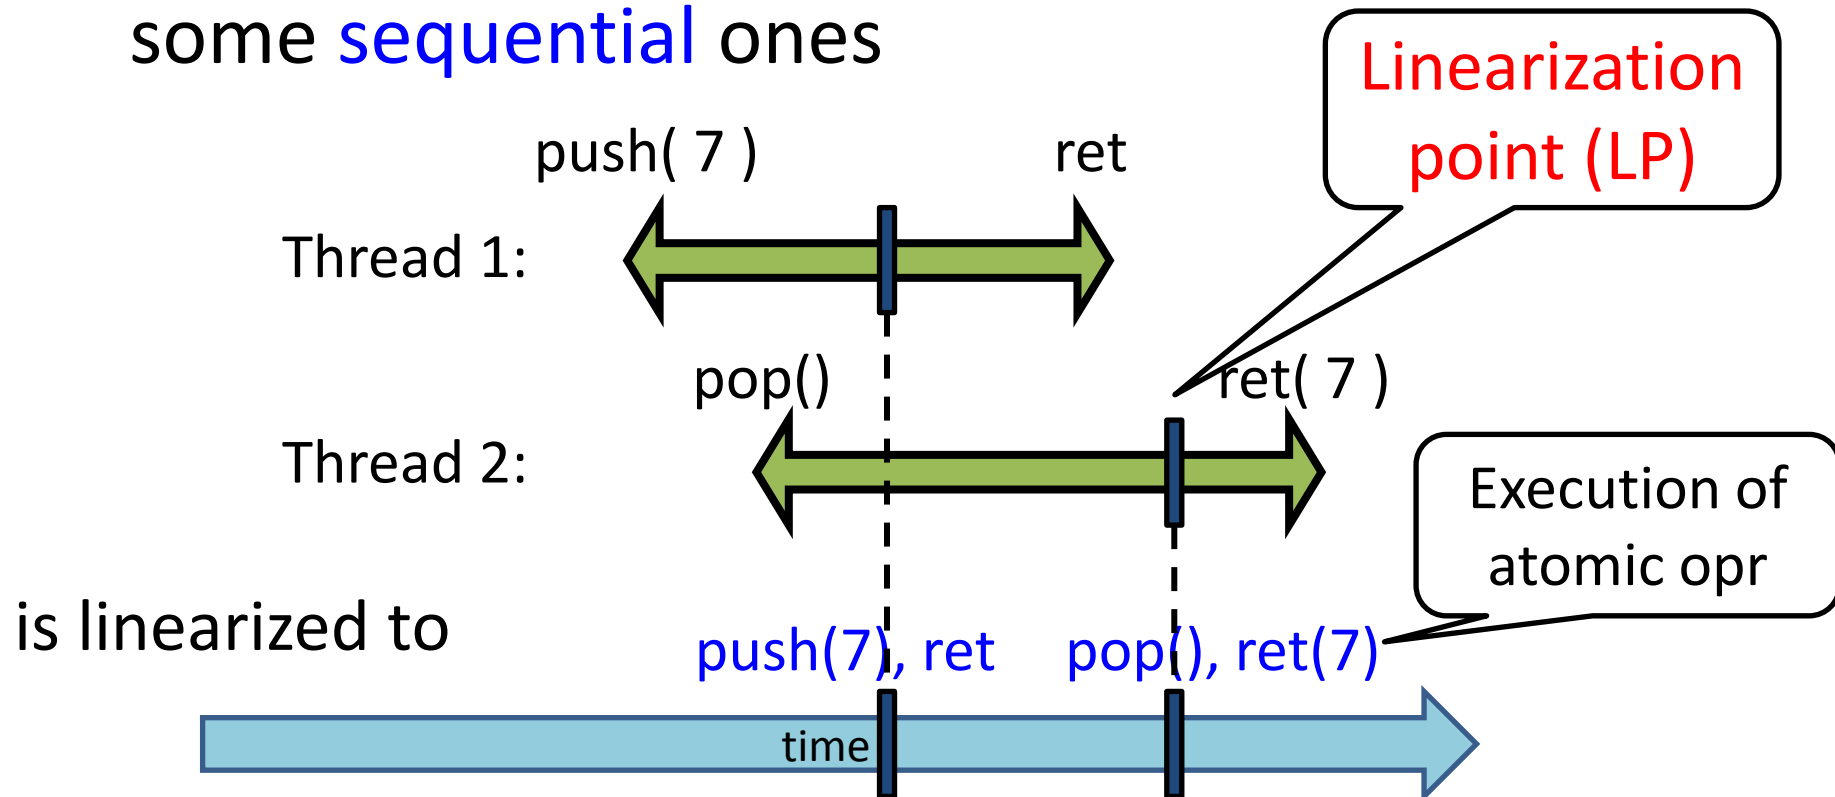
*How to specify/prove correctness?*

```
push(int v):  
1  int d=0, x, t;  
➔ 2  x = new Node();  
3  x.data = v;
```

```
4  while(d==0){  
5      <t = top>;  
6      x.next = t;  
7      d = cas(&top, t, x);  
8  }
```

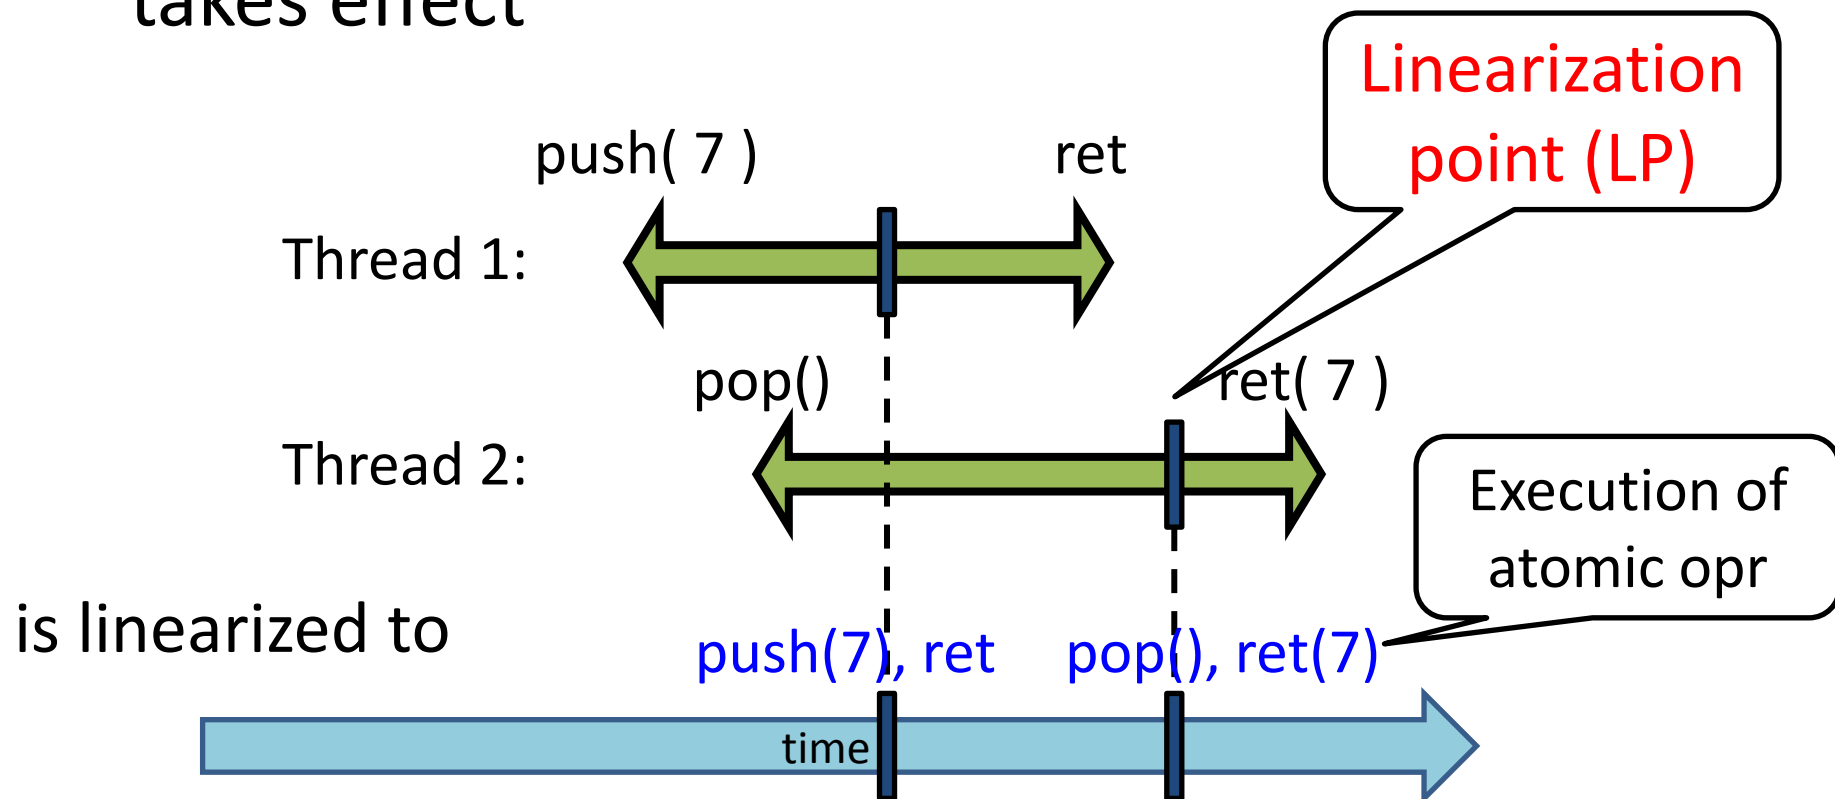
# Linearizability [Herlihy & Wing'90]


- Standard correctness criterion
- All **concurrent** executions are “equivalent” to some **sequential** ones



# LP Method to Prove Linearizability

- Locate **LP** in **impl code** **O**
- Show it is the single point where the method takes effect





## Treiber's Stack

**push(v):**

1 local d:=0, x, t;

2 x := new Node(v);

3 while(d=0){

4 t := top;

5 x.next := t;

6 **d := cas(&top, t, x);** LP

7 }

Not update the shared list

Line 6: the only command that changes the list


# Problems with LP Method

- Informal
  - Mostly folklore theorem
- Difficult to locate LP
  - LP cannot be statically located – non-fixed LP
  - Common in many wait-free algorithms


# Example: HSY Stack

**T1:**  
**push(v)**

**Elimination  
Array**



**T2:**  
**pop()**



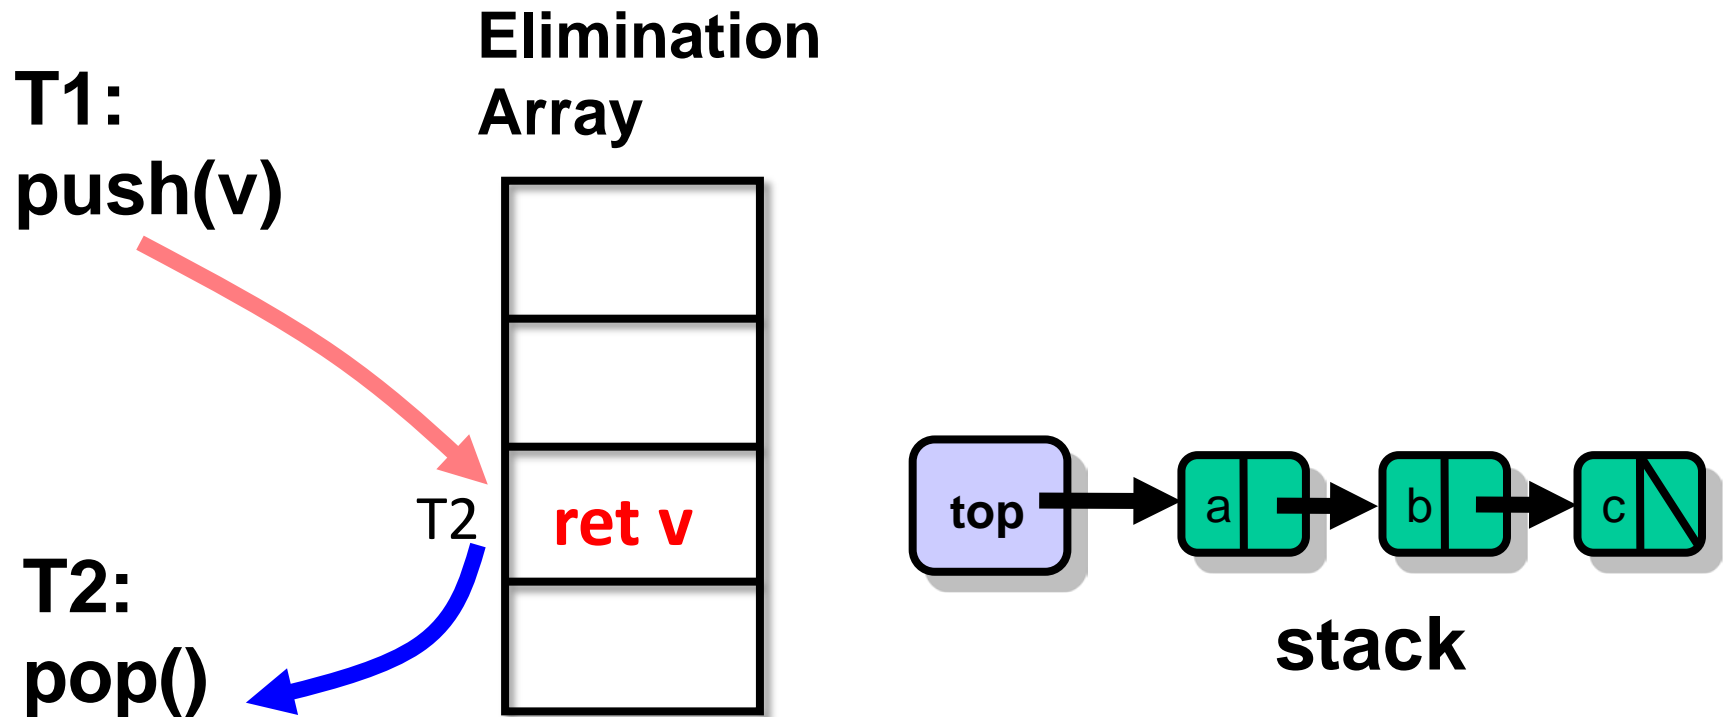
**stack**

Find T2 is doing p

**Thread  
descriptor**




# Example: HSY Stack



T1 finishes not only its own opr, but also T2's

# Coordination Pattern in Wait-Free Alg.



T1 interrupts T2, and helps T2 to finish its pending opr.

When T2 comes, its job is done.

*Difficult to find LP of T2's opr!*

# Problems with LP Method

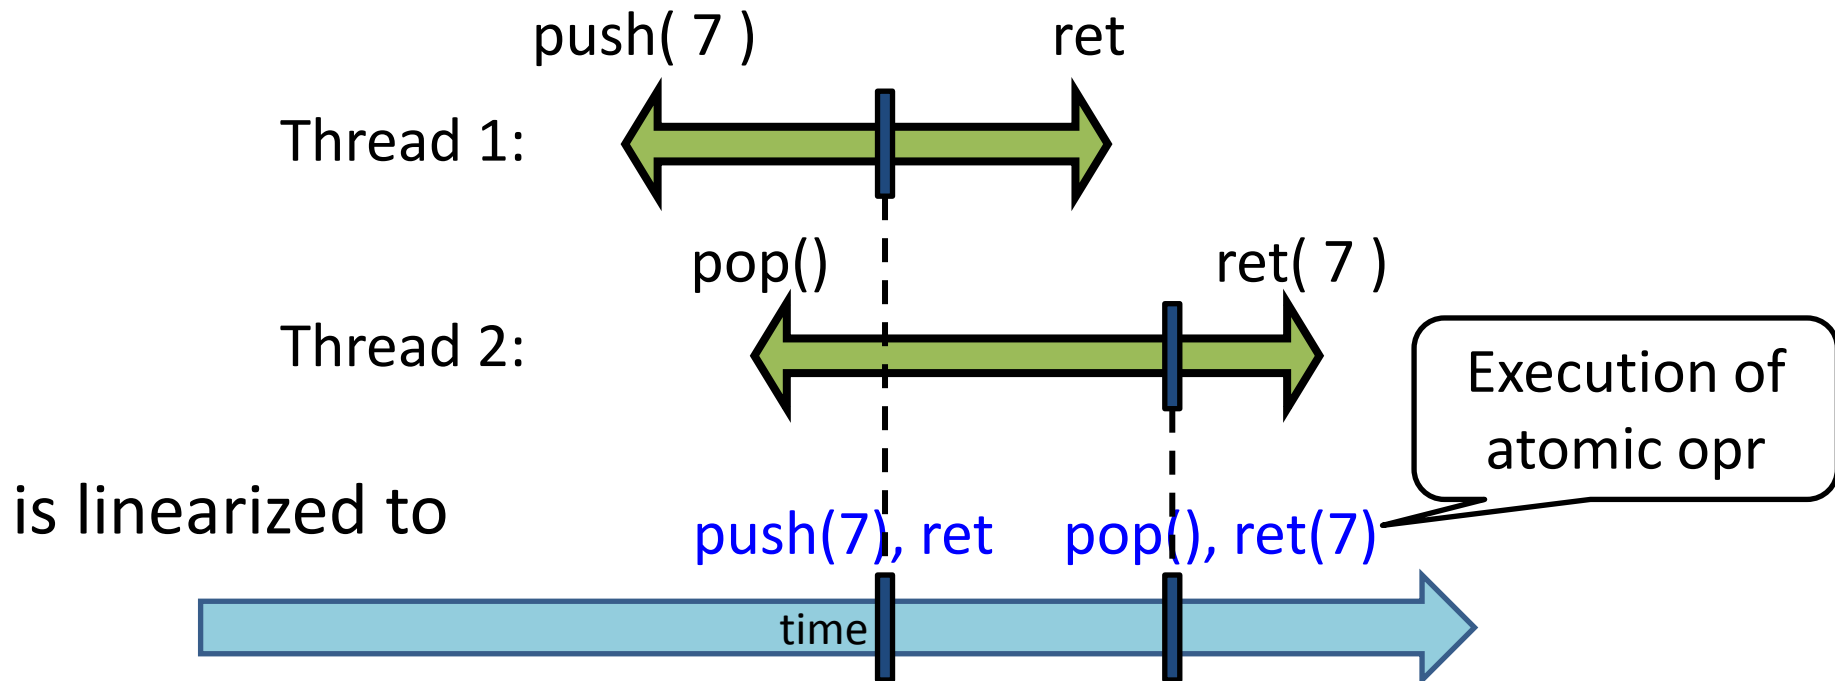
- Informal
  - Mostly folklore theorem
- Difficult to locate LP
  - LP cannot be statically located – non-fixed LP
  - Common in many wait-free algorithms

# This Talk

- Simulation-based proof for linearizability
  - Supports compositional verification
  - Formally justifies the folklore LP method
- Extending it for objects with non-fixed LP
  - HSY elimination stack, lazy set , etc.

# Linearizability and Program Refinement


**Observation:** Linearizable fine-grained impl.  
has the same effect as atomic operations




# Linearizability and Program Refinement

Observation: Linearizable fine-grained impl.  
has the same effect as atomic operations

***Reduce linearizability to program equivalence:***





## Treiber's Stack

**push(v):**

```
1 local d:=0, x, t;  
2 x := new Node(v);  
3 while(d=0){  
4   t := top;  
5   x.next := t;  
6   d := cas(&top, t, x);  
7 }
```

object impl **O**

**Abstract representation**

stk: n1 :: n2 :: ... :: nk


**push(v):**

**<stk := v::stk>;**

**Atomic operation**

atomic opr as spec **A**


# Simulation for Refinement Proof




Whatever behavior produced by low-level could be produced by high-level.



# Customized for Linearizability





**push(v):**

1 local d:=0, x, t;

2 x := new Node(v);

3 while(d=0){

4 t := top;

5 x.next := t;

6 d := cas(&top, t, x);

7 }

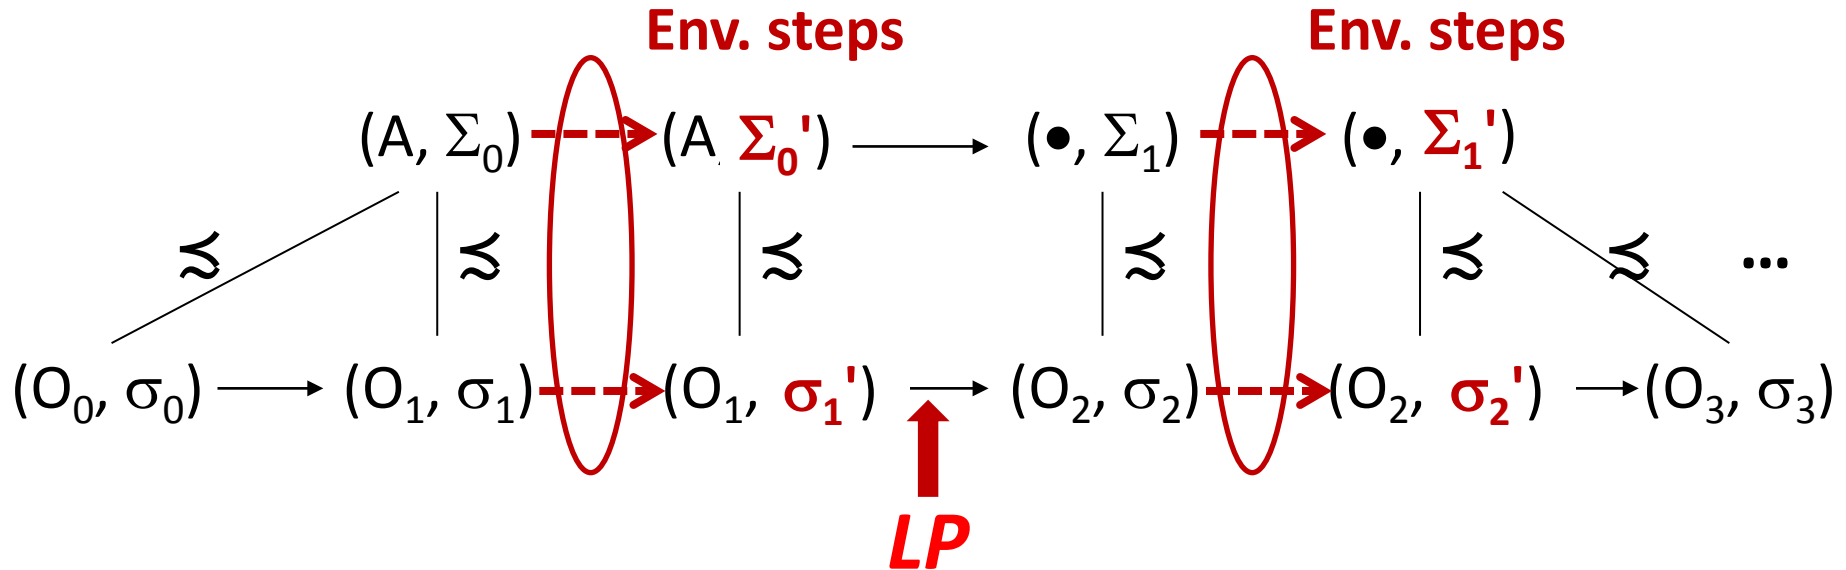
However, needs to consider env., otherwise unsound.



**<stk := v::stk>;**



# Simulation with Env.

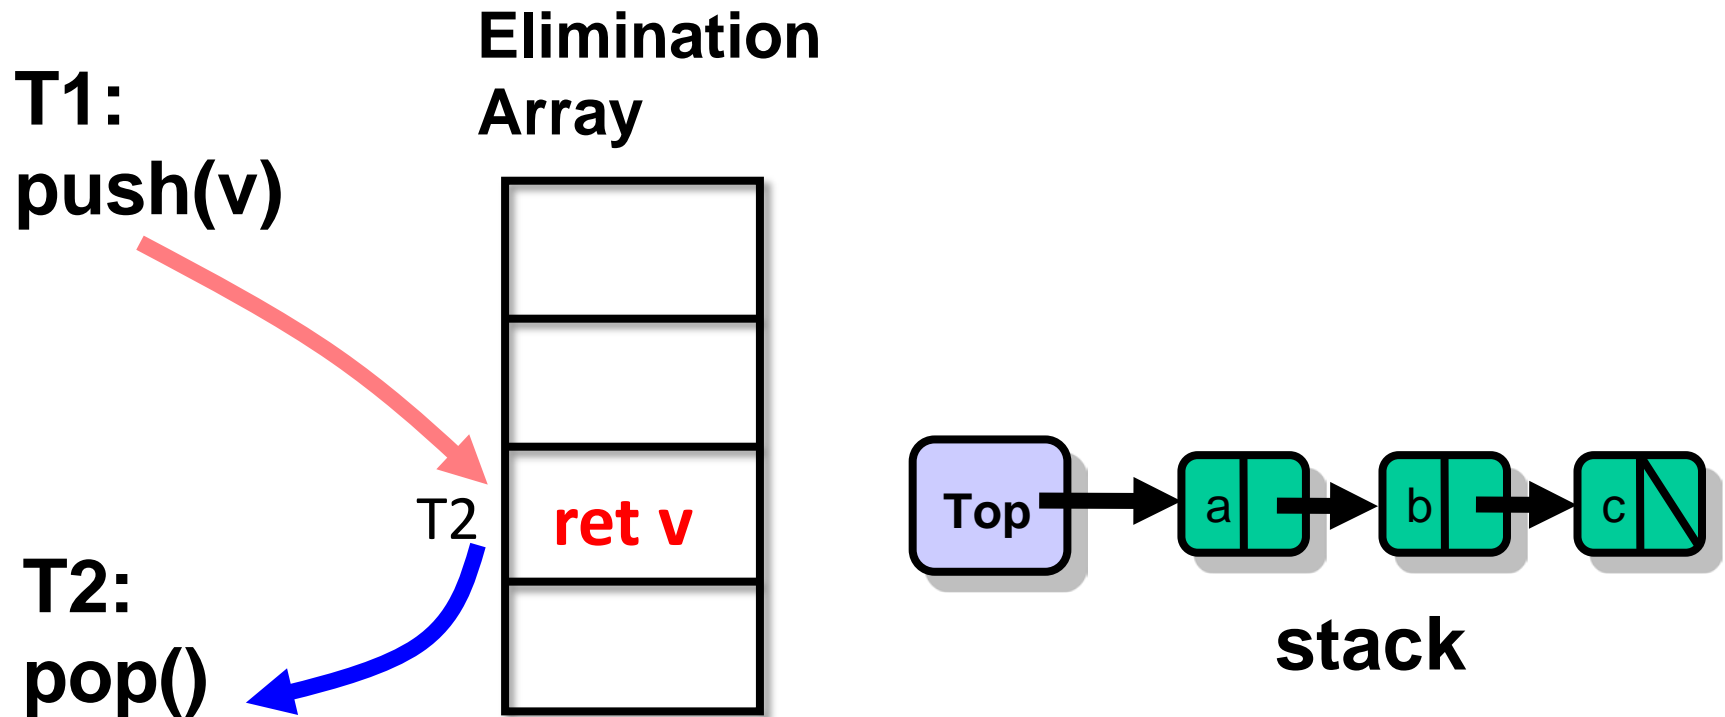


Needs to ensure the env. steps do not break simulation  
(environment doesn't do bad things)

# Simulation with Env.


- RGSim: Rely-Guarantee based simulation [Liang et al'12]
  - General method for concurrent program refinement
  - Adapted for linearizability proof
    - Formally justifies the LP method
    - Used to verify Treiber's stack
- Needs to know statically LP point
  - Does not support non-fixed LP as in HSY stack

# HSY Stack Implementation



T1 finishes not only its own opr, but also T2's

# HSY Stack Implementation



What's the problem?


The LP of pop may not even be inside the method body.

LP corresponds to env.  
Step, not supported in  
previous sim.

T1 finishes not only its own opr, but also T2's

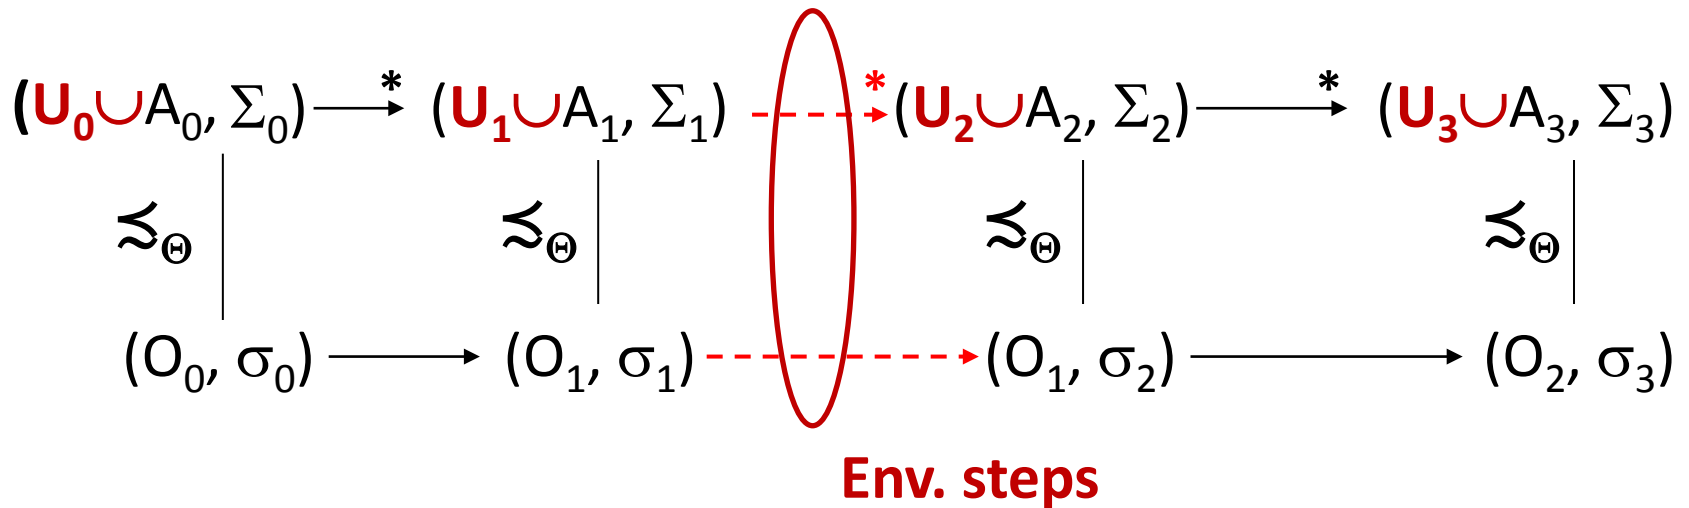
# Our Solution

- Parameterize RGSim with **pending threads** that might be helped by others



- T1's step may fulfill opr of T1 or threads in U
- T2 checks U to see if its opr has been done (by env)

# Our New Simulation $O \lesssim_{\Theta} S$




$U$  maps a thread ID  $t$  to a pending opr.  $A$  or  $\bullet$

$$U_i \cup A_i = \Theta(\sigma_i)$$


Support non-fixed LPs: thread's LP can be in either  $\longrightarrow$  or  $\xrightarrow{\text{dashed}}$





$\mathbf{U}(\mathbf{T2}) = \text{pop}$

$\mathbf{U}'(\mathbf{T2}) = \bullet$



Soundness:

If  $O \lesssim_{\Theta} A$  for some  $\Theta$ , then  $O \leq_{\text{lin}} A$

What's more:

Hoare-style syntactic logic for lin.

<http://home.ustc.edu.cn/~lhj1018/lin.pdf>

# Conclusion

- A new simulation  $O \preceq_{\text{H}} A$  for linearizability proof
  - Sim  $\rightarrow$  refinement  $\rightarrow$  linearizability
  - Supports non-fixed LP
    - HSY elimination-based stack, lazy set, etc
    - First refinement-based proof for HSY elimination-based stack
  - A program logic for syntactic verification
- Another dimension of complexity
  - LP depends on future
  - May need backward simulation (leave as future work)

Thank you!