



Challenges of self-adaptive software

*the fading boundary between
development time and run time*

Internetware 2012 & ISHCS 2012, Qingdao (China)

Carlo Ghezzi
Politecnico di Milano
Deep-SE Group @ DEI-PoliMI



The vision

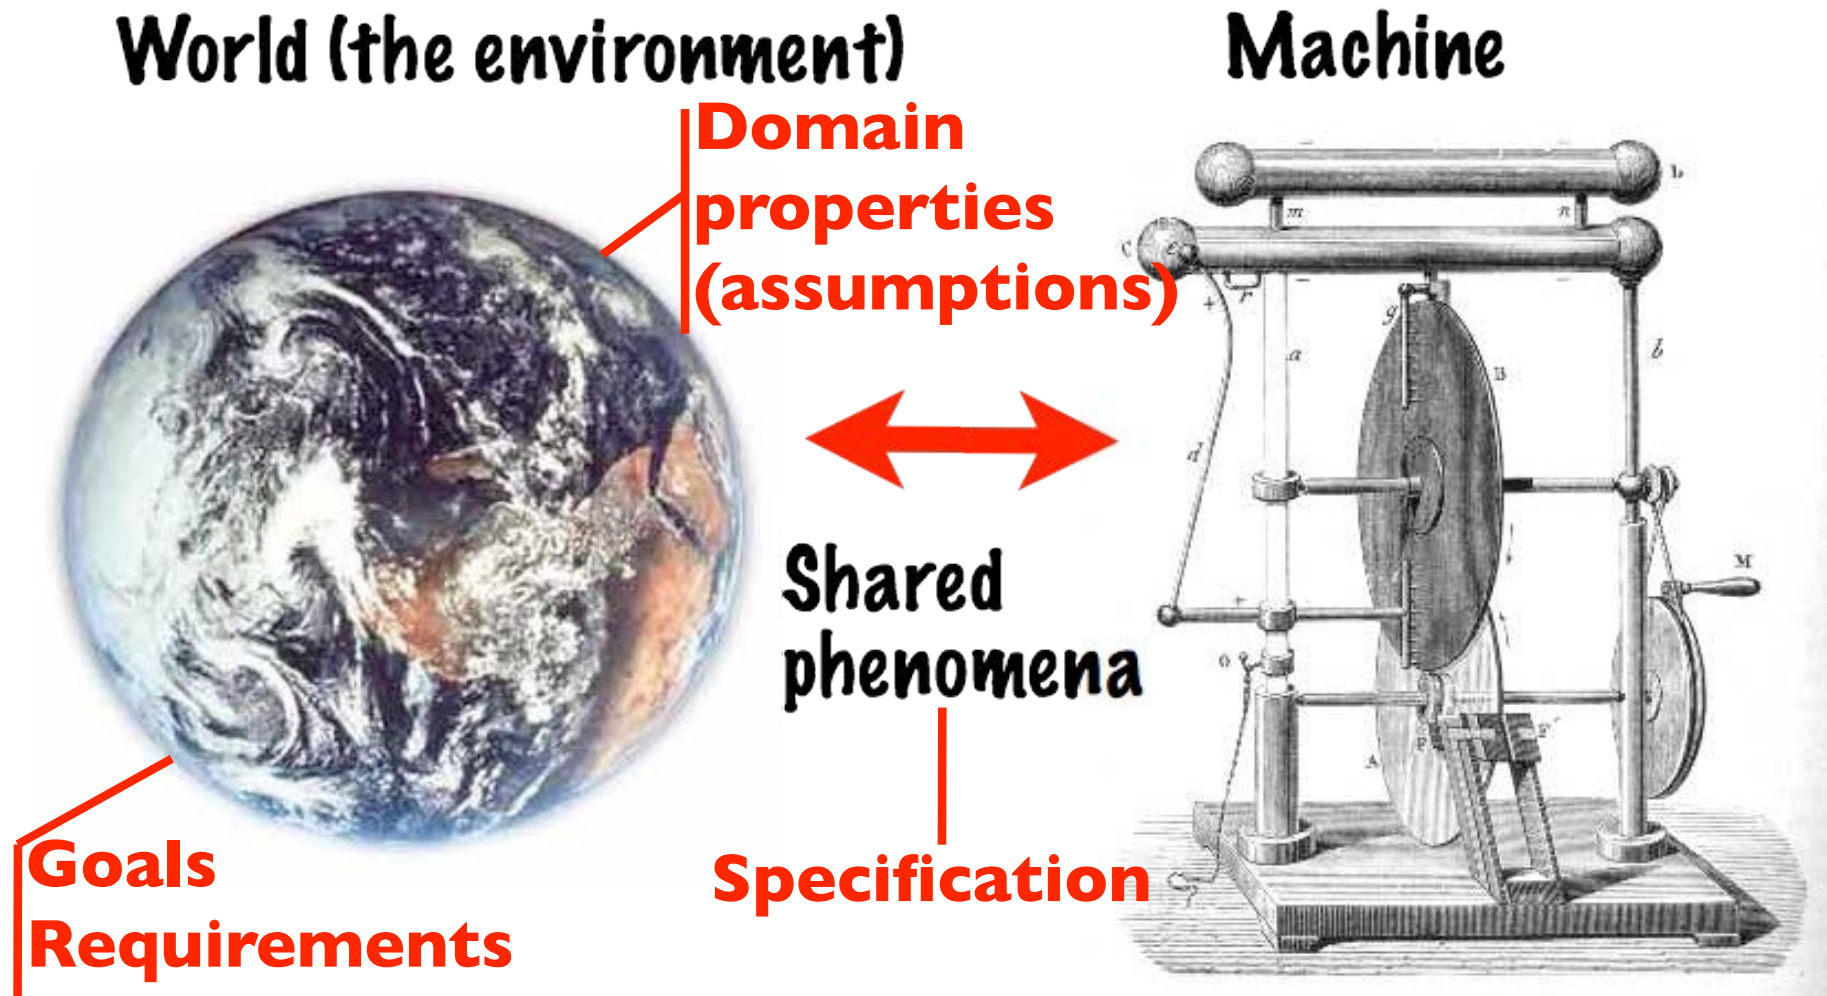
- **World fully populated by computationally rich devices offering services (disappearing computer)**
 - appliances, sensors/actuators, ... “things”
- **Cyber-physical systems**
- **Mobility**
- **Situation-aware computing**
 - new “services” built dynamically in a situation-dependent manner
- **Continuously running systems**
 - need to evolve while they offer service



The challenge

- Change and flexibility adversary of dependability
- Can they be reconciled through sound design methods?

The *machine* and the world





Dependability arguments

- Assume that a rigorous (formal) representation is given for
 - R = requirements
 - S = specification
 - D = domain assumptions

if S and D are all satisfied and consistent, it is necessary to prove

$$- S, D \models R$$



Change comes into play

- Changes in **goals/requirements**
- Changes in **domain assumptions**
 - **Usage context**
 - request profiles
 - **Physical context**
 - space, time, ...
 - **Computational context**
 - external components/services (*multiple ownership*)
 - systems increasingly built out of parts that are developed, maintained, and even operated by independent parties
 - no single stakeholder oversees all parts, which may change independently
 - yet by assembling the whole one commits to achieving certain goals



Changes may affect dependability

- Changes may concern
 - R **evolution**
 - D **adaptation** *here I focus on D*
- We can decompose D into D_f and D_c
 - D_f is the fixed/stable part
 - D_c is the changeable part

We need to **detect changes to D_c** **change detection**
and **make changes to S** to keep satisfying R
(self) adaptation




Change revisited


- Change recognized as a crucial problem since the 1970's (see work by M. Lehman)
- Traditionally managed off-line: **software maintenance**
- What is new here
 - the unprecedented degree of change
 - the request that software responds to changes while the system is running (continuously running systems), possibly in a **self-managed** manner



A paradigm change

- Conventional separation between development time and run time is blurring
- Models + requirements need to be kept + updated at run time
- Continuous verification must be performed to detect the need for adaptation





Zoom-in

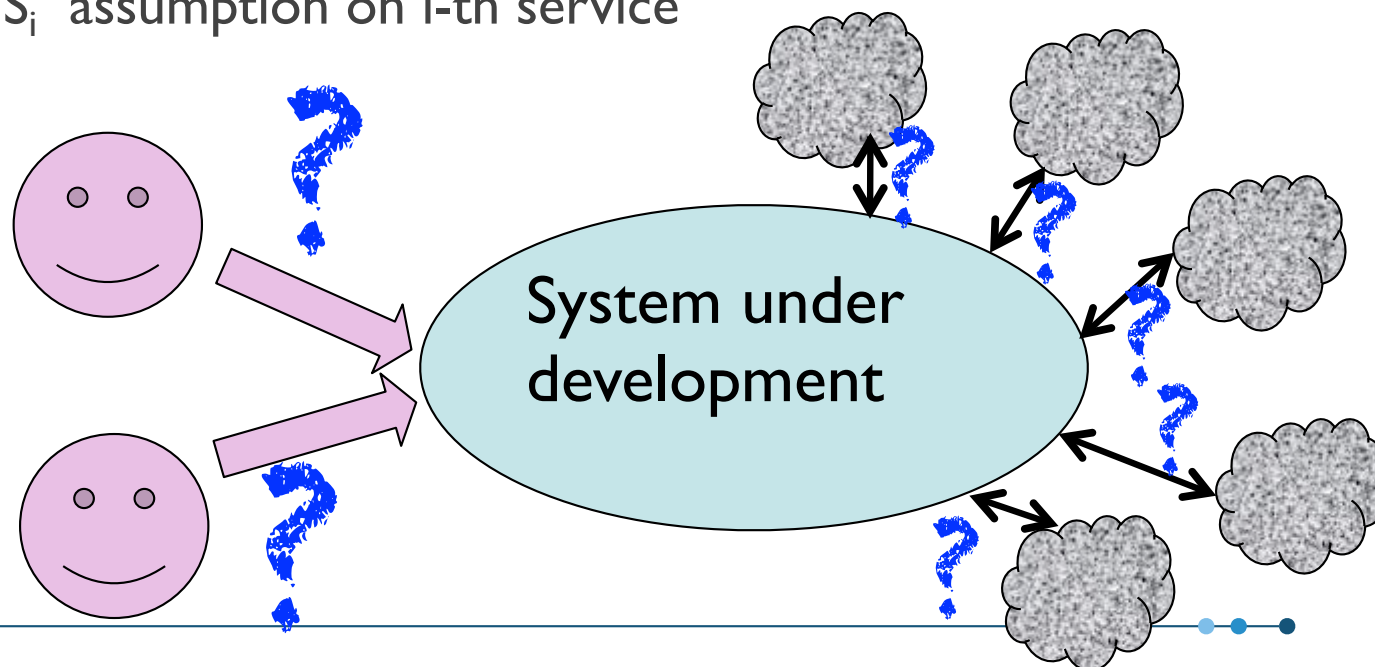
A framework for (self) adaptation

- I. Epifani, C. Ghezzi, R. Mirandola, G. Tamburrelli, "Model Evolution by Run-Time Parameter Adaptation", ICSE 2009
- C. Ghezzi, G. Tamburrelli, "Reasoning on Non Functional Requirements for Integrated Services", RE 2009
- I. Epifani, C. Ghezzi, G. Tamburrelli, "Change-Point Detection for Black-Box Services", FSE 2010

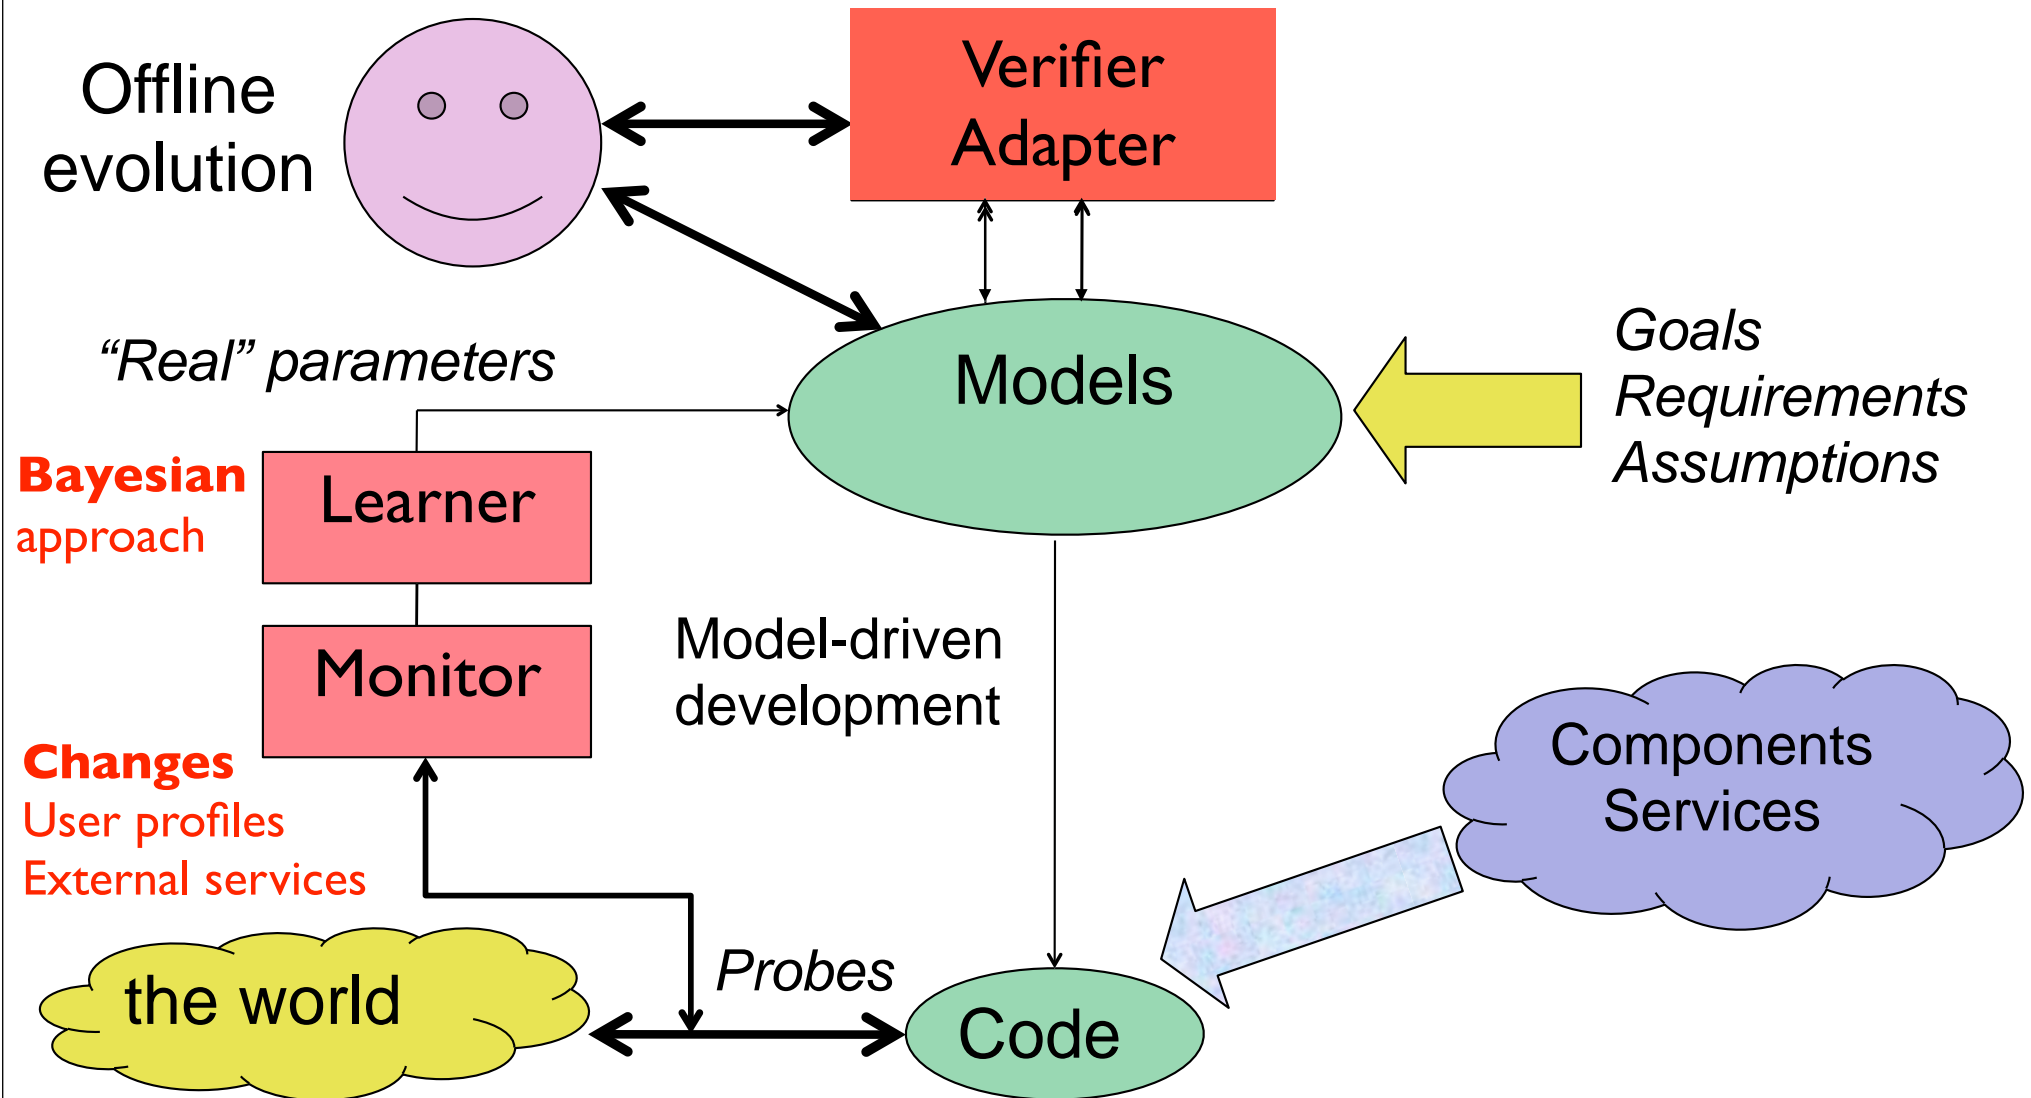
Specific focus

- **Non-functional** requirements
 - reliability, performance, energy consumption, cost, ...
- Quantitatively stated in **probabilistic** terms
- D_c decomposed into D_u, D_s
 - D_u = usage profile
 - $D_s = S_1 \wedge \dots \wedge S_n$ S_i assumption on i-th service

Hard to estimate at design time + very likely to change at run time



Our approach in a nutshell





Models

- Different models provide different **viewpoints** from which a system can be analyzed
- Focus on **non-functional** properties and quantitative ways to deal with uncertainty
- Use of **Markov models**
 - DTMCs for reliability
 - CTMCs for performance
 - Reward DTMCs for energy/cost

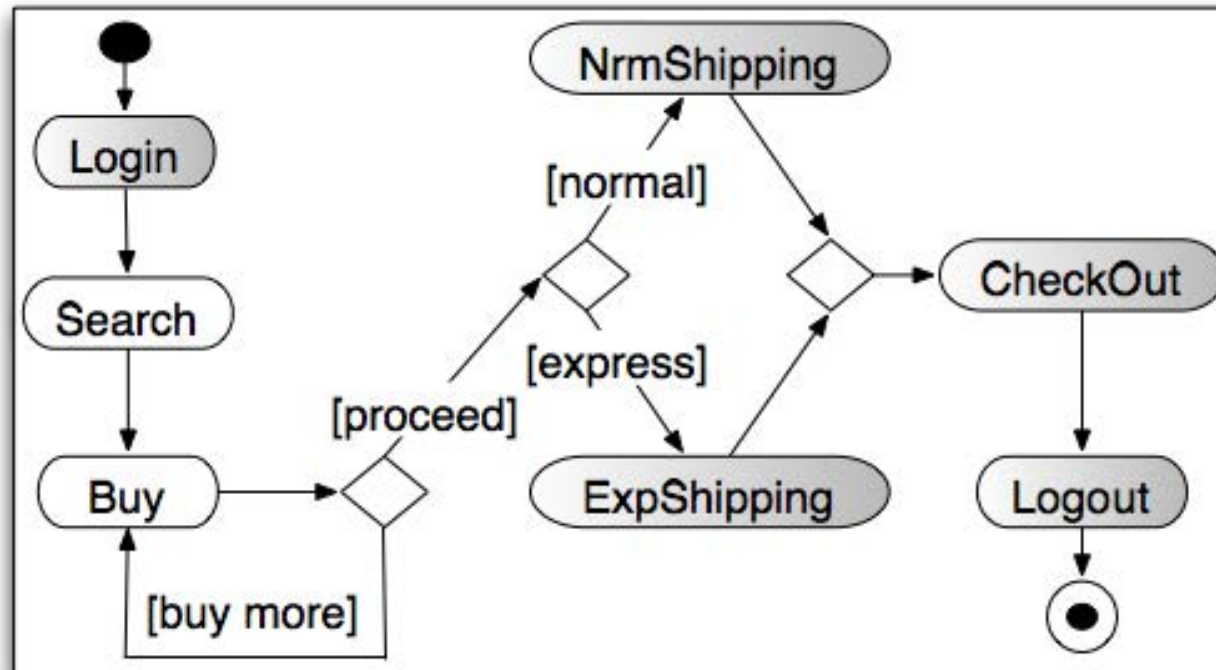


Properties and verification (the case of *reliability*)

- PCTL (probabilistic extension of CTL) provides the necessary expressive power
 - most reliability specifications can be stated as reachability properties
 - $P_{>0.8} [\Diamond(\text{system state} = \text{success})]$
- excellent tools exist to perform verification via model checking
 - PRISM (Kwiatkowska et al.)
 - <http://www.prismmodelchecker.org/>
 - MRMC (Katoen et al.)
 - <http://www.mrmc-tool.org/trac/>



The approach in in action: e-commerce service composition



Users
classified as BigSpender
or SmallSpender based on their
usage profile.

3 probabilistic requirements:

R1: "Probability of success is > 0.8 "

R2: "Probability of a ExpShipping failure for a user recognized as BigSpender < 0.035 "

R3: "Probability of an authentication failure is less then < 0.06 "



Assumptions

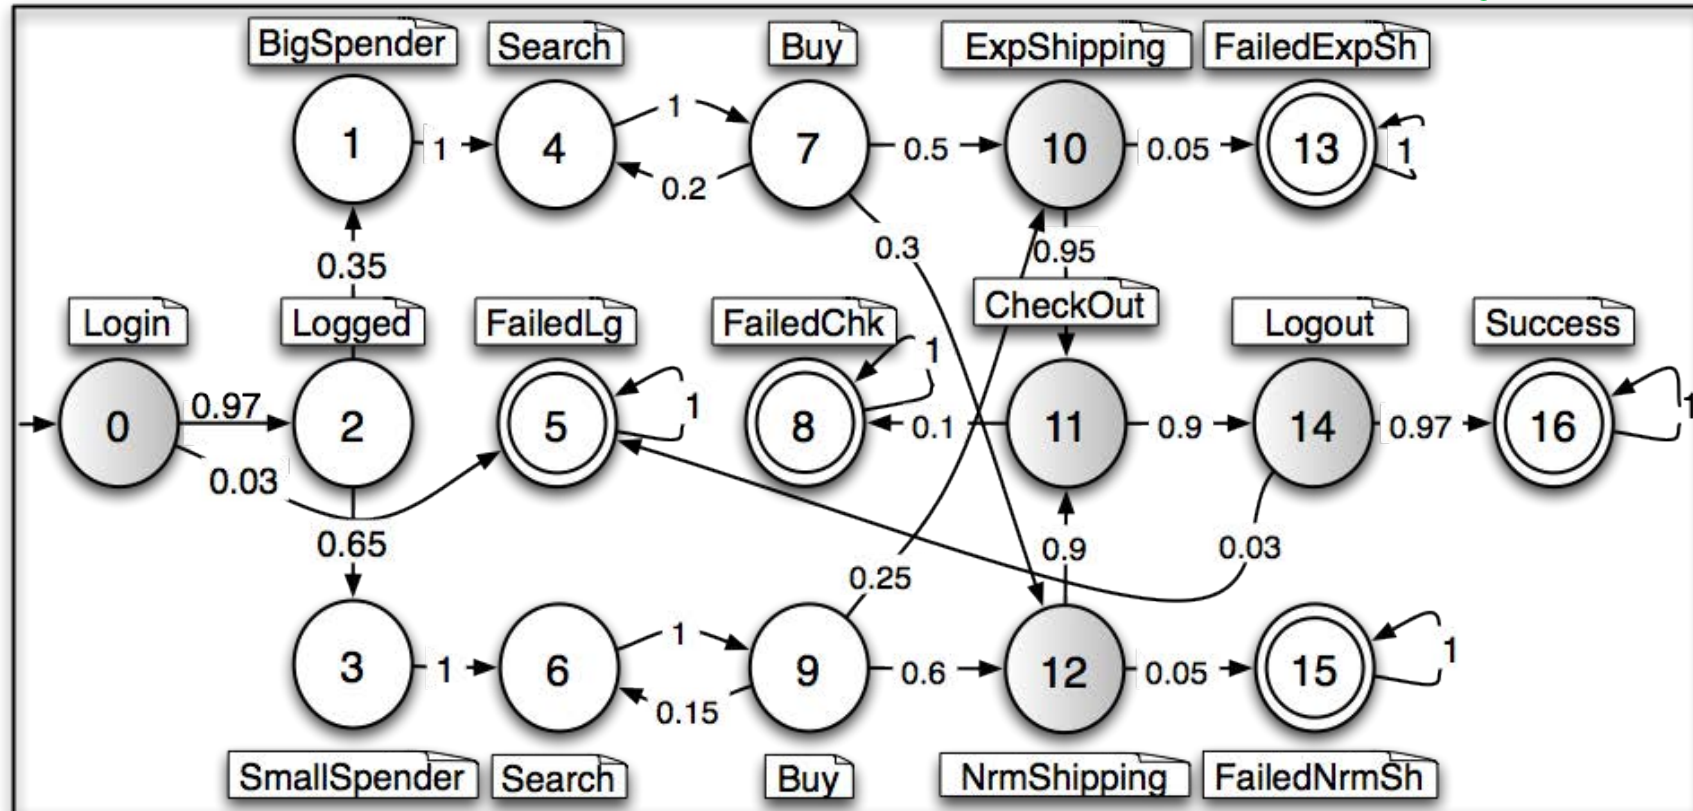
User profile domain knowledge

$D_{u,n}$	Description	Value
$D_{u,1}$	$P(\text{User is a BS})$	0.35
$D_{u,2}$	$P(\text{BS chooses express shipping})$	0.5
$D_{u,3}$	$P(\text{SS chooses express shipping})$	0.25
$D_{u,4}$	$P(\text{BS searches again after a buy operation})$	0.2
$D_{u,5}$	$P(\text{SS searches again after a buy operation})$	0.15

External service assumptions (reliability)

$D_{s,n}$	Description	Value
$D_{s,1}$	$P(\text{Login})$	0.03
$D_{s,2}$	$P(\text{Logout})$	0.03
$D_{s,3}$	$P(\text{NrmShipping})$	0.05
$D_{s,4}$	$P(\text{ExpShipping})$	0.05
$D_{s,5}$	$P(\text{CheckOut})$	0.1

DTMC model



Property check via model checking

R1: "Probability of success is > 0.8 " **0.84**

R2: "Probability of a ExpShipping failure for a user recognized as BigSpender < 0.035 " **0.031**

R3: "Probability of an authentication failure is less then < 0.06 " **0.056**



What happens at run time?

- We monitor the actual behavior
- A statistical (Bayesian) approach estimates the updated DTMC matrix (posterior) given run time traces and prior transitions
- Boils down to the following updating rule

$$m_{i,j}^{(N_i)} = \frac{c_i^{(0)}}{c_i^{(0)} + N_i} \times m_{i,j}^{(0)} + \frac{N_i}{c_i^{(0)} + N_i} \times \frac{\sum_{h=1}^d N_{i,j}^{(h)}}{N_i}$$


A-priori Knowledge A-posteriori Knowledge




Faults and failures

- **Fault**
 - Machine or environment do not behave as expected
- **Failure**
 - Experienced violation of requirement
- Assume that an *environment* fault is detected
Three cases are possible
 - All Reqs still valid
 - **OK, but my signal contract violation**
 - Some Req violated + violation experienced in real world
 - **Failure detection**
 - Some Req violated, but violation not experience yet
 - **Failure prediction**

Predicted vs. detected failure



Suppose that execution traces that lead to updating the failure probability of ExpShipping are those involving small spenders
 R2: Probability of a ExpShipping failure for a user recognized as BigSpender < 0.035



Zoom-in

Run-time efficient model checking

A. Filieri, C. Ghezzi, G. Tamburrelli, “Run-Time Efficient Probabilistic Model Checking”, ICSE 2011

A. Filieri, C. Ghezzi, G. Tamburrelli, “A formal approach to adaptive software: continuous assurance of non-functional requirements”, Formal Aspects of Computing, vol. 4, n. 2, 2012

C. Ghezzi, "Evolution, Adaptation, and the Quest for Incrementality", in Large-Scale Complex IT Systems. Development, Operation and Management, 17th Monterey Workshop 2012, Oxford, UK, March 19-21, 2012, LNCS, Springer 7539

Rethinking run-time environments



- Traditionally software engineering has been mostly concerned with development time
- The result is **code** that simply needs to be **run**

(Self-)adaptive software requires much more

- *must be able to reason at run time about itself and the environment*

- ✓ *models*

- ✓ *goals and requirements*

- ✓ *strategies*

must be available at runtime



Run-time agility, incrementality

- Agility taken to extremes
 - time boundaries shrink
 - ✓ constrained by real-time requirements
- Verification approaches must be re-visited
 - they must be **incremental**

Given S system (model), P property to verify for S

Change = new pair S' , P'

Incremental verification reuses part of the proof of S against P to verify S' against P'



How to make verification incremental

Incrementality by encapsulation

- Grounded on seminal work of D. Parnas (1972)
 - Design for change
 - ✓ changes must be anticipated and encapsulated within *modules*
 - ✓ interface vs implementation
 - ✓ interfaces formalized via contracts (B. Meyer)
- Known as *assume-guarantee* when contextualized to verification (C. Jones)



Assume-guarantee

- Show that module $M1$ guarantees property $P1$ assuming that module $M2$ delivers property $P2$, and vice versa
- Then claim that the system composed of $M1 \parallel M2$ guarantees $P1$ and $P2$ unconditionally
 - these arguments support *compositional reasoning*
- Approach works if changes do not percolate through the module's interface, affecting contract
 - effect of change *encapsulated* within the boundaries predicted at design time



How to make verification incremental

Incrementality by parameterization

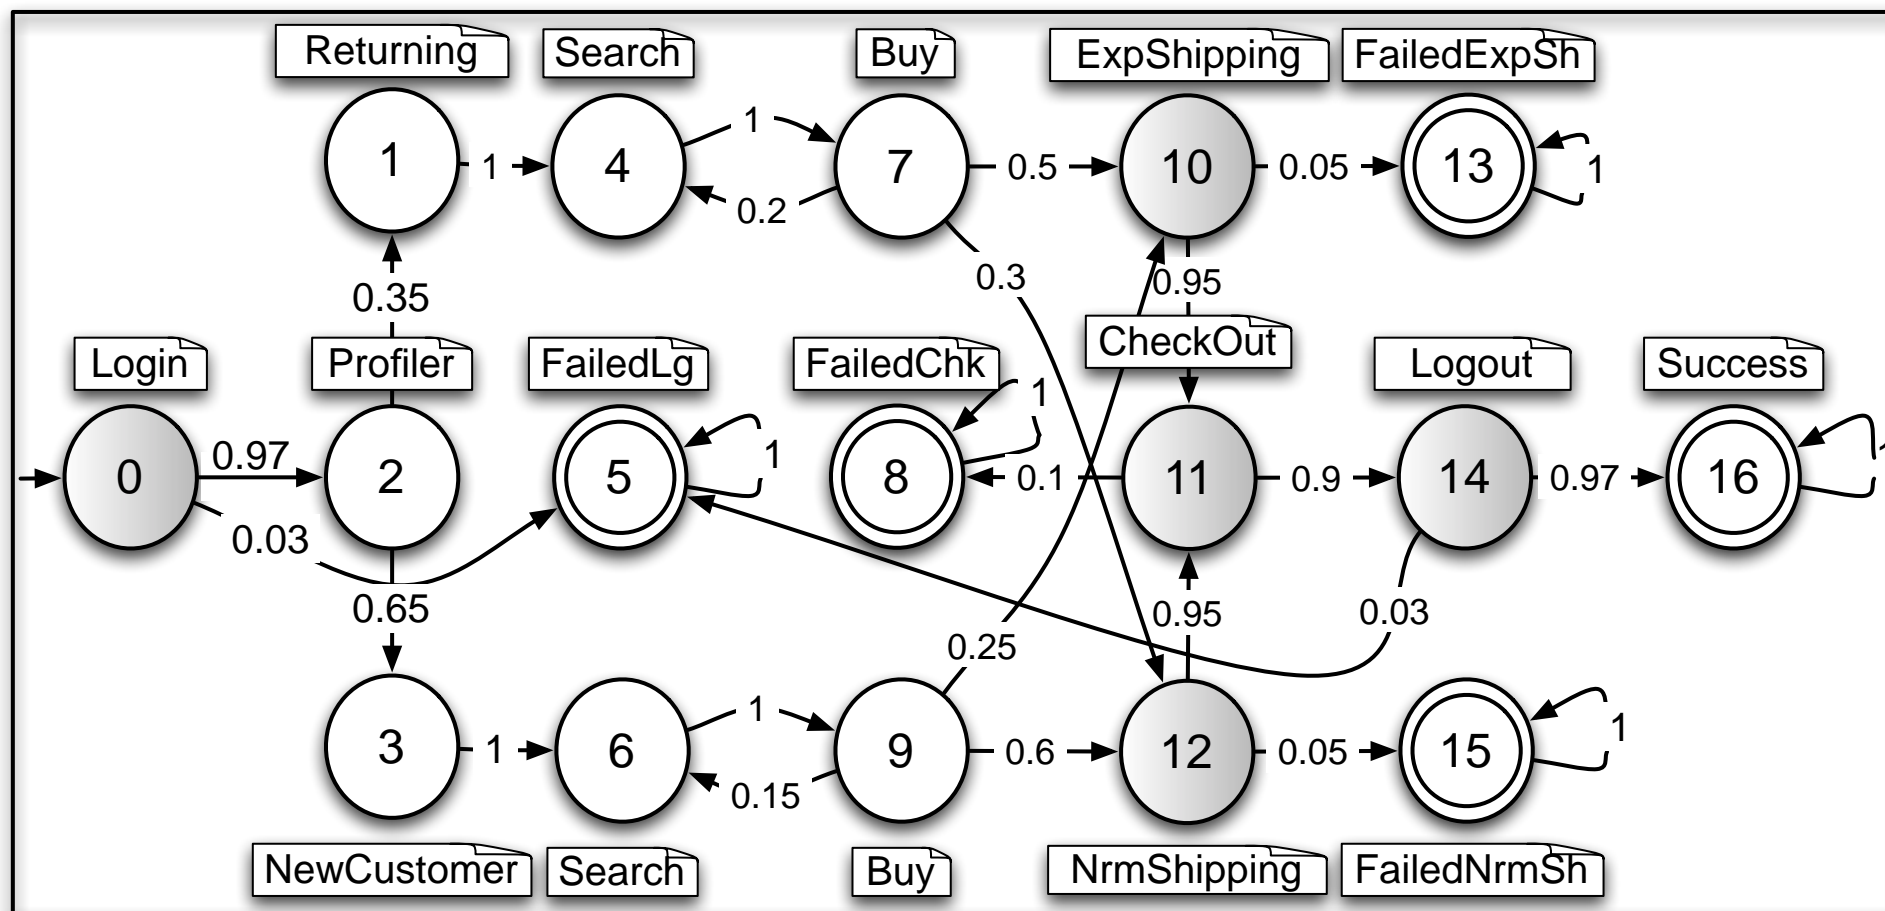
- Requires anticipation of changes, which become parameters
- Does not require modular reasoning
- Still requires identification of elementary sources of change
- Inspired by the concept of partial evaluation (Ershov 1977)

Let P be a program $P : I \rightarrow O$

Suppose I can be partitioned into I_s and I_d , where I_s set of input data known statically, before runtime

Partial evaluation transforms P into an equivalent *residual* program $P' : I_d \rightarrow O$ from by precomputing static input before runtime

An example

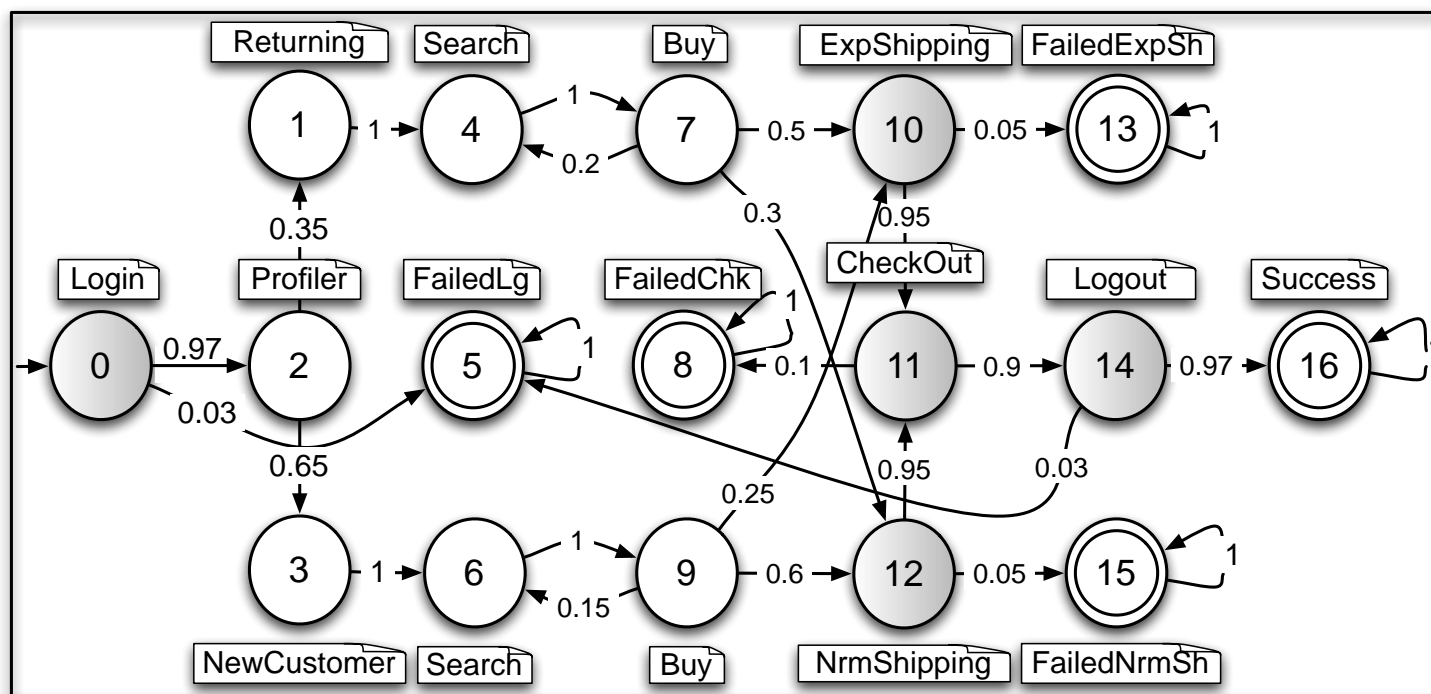


Requirement: **$F_{>0.8} [s = 16]$**

An example (continues)

Satisfaction of requirement **$F_{>0.8} [s = 16]$** can be checked at design time, but at run-time, e.g. user profiles may change

We can treat them as **variables** and compute at design time a parametric verification formula which is then evaluated at run time



How to make verification incremental

Syntax-driven incrementality

- Assumes artifact to analyze with a syntactic structure expressible as a formal grammar
- Verification is expressed via attributes (à la Knuth)
- Changes can be of any kind






Intuition

Syntax-driven incrementality

- Incremental parsing strategy finds boundary for artifact re-analysis
- Knuth proved that attributes can be only **synthesized** (computed bottom-up) and thus only need to be recomputed for the changed portion + propagated to the root node

Incremental parsing: intuition



- Assume w is the modified portion
- Ideally, re-analyze only a sub-tree “covering” w , rooted in $\langle N \rangle$, and “plug-it-in” the unmodified portion of tree
- The technique works if the sub-tree is small, and complexity of re-analysis is the same as complexity of “main” algorithm



Incremental parsing: past and new results

- Past work on “mainstream” LR grammars
 - C. Ghezzi and D. Mandrioli, Incremental parsing, ACM Trans. Program. Lang. Systems, 1979
 - ✓ Saves the maximum possible portion of the syntax tree, but the re-analyzed portion can still be large in certain cases
- Recent work resurrected Floyd’s operator precedence grammars
 - R.W. Floyd. Syntactic analysis and operator precedence, Journal of the ACM, 1963
 - S. Crespi Reghizzi and D. Mandrioli. Operator-precedence and the visibly pushdown property, J. Comput. Syst. Sci., to appear.
 - Floyd’s grammars cannot generate all deterministic CF languages
 - but in practice any programming language can be described by a Floyd grammar
 - parsing can be started from any arbitrary point of the artifact to be analyzed




Initial validation of the approach

- Case I: reliability (QoS) analysis of composite workflows
 - a (BPEL) workflow integrates external Web services having given reliability and we wish to assess reliability of composition
 - if reliability of an external service changes, does our property about reliability of composition change?
 - ✓ our previous work framed this into probabilistic model checking
 - ✓ here we can deal with unrestricted changes, also in the workflow in a very efficient way



Initial validation of the approach

- Case 2: reachability analysis as supported by program model checking
 - given a program and a safety property, is there an execution of the program that leads to a violation of the property?
 - if the program changes, how does our property change?
 - ✓ similar problem faced by Henzinger et al.
 - T.A. Henzinger, R. Jhala, R. Majumdar, and M.A. Sanvido. Extreme model checking. In Verification: Theory and Practice, volume 2772 of LNCS, 2004.



Zoom-in

Control-theory based self adaptation

[ASE 2011] A. Filieri, C. Ghezzi, A. Leva, M. Maggio,
Self-Adaptive Software Meets Control Theory: A Preliminary
Approach Supporting Reliability



Goal

- Tune software through its model via feedback control loop
- Formally prove controller's capabilities (error reduction, convergence, ...)



Conclusions and future work

- (Self-)adaptation is needed
- It requires a paradigm shift
- Run-time environments must become semantically rich
- Run-time **reasoning** must be supported, not just execution
- Continuous change and the quest for incrementality
- Benefits can be achieved by applying methods from control theory



Thanks to the group





Acknowledgements

This work is supported by and Advanced Grant of the
European Research Council
Programme IDEAS-ERC, Project 227977---SMScom