

Microsoft Research Software Analytics Group

Context-Sensitive Delta Inference for Identifying Workload-Dependent Performance Bottlenecks

Xusheng Xiao¹, Shi Han², Dongmei Zhang², **Tao Xie**^{1,3}

¹North Carolina State University

²Microsoft Research Asia

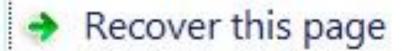
³University of Illinois at Urbana-Champaign

¹xxiao2@ncsu.edu, ²{shihan, dongmeiz}@microsoft.com, ^{1,3}taoxie@illinois.edu

Windows Internet Explorer

A webpage is not responding on the following website: facebook.com

You can wait for the webpage to respond, or choose one of the following options:



Close this page



See details





This program is not responding.

To return to Windows and check the status of the program, click Cancel.

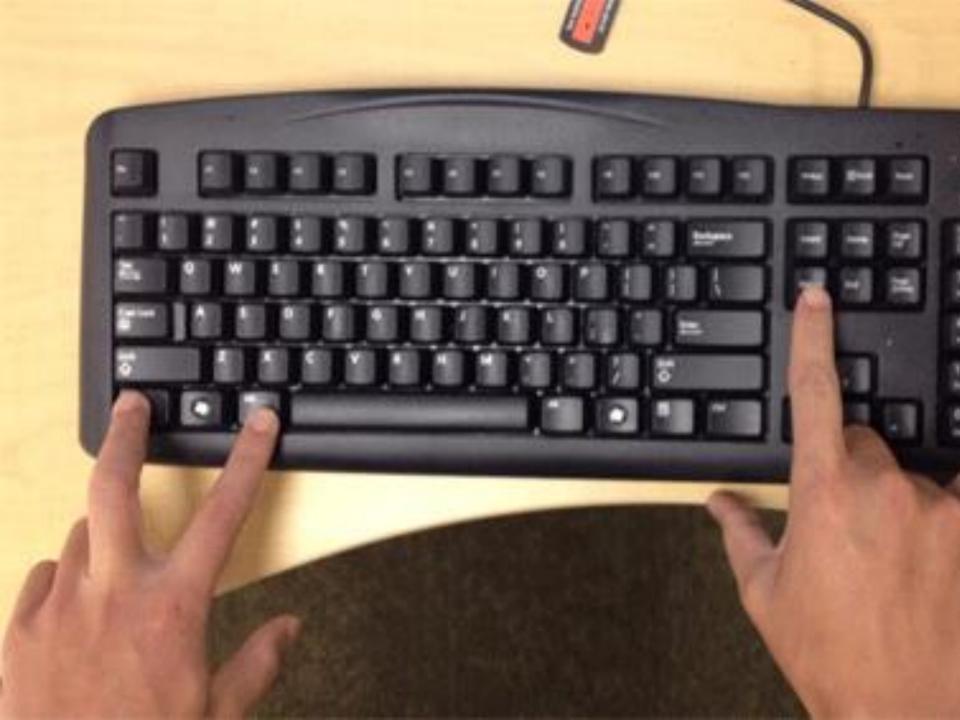
If you choose to end the program immediately, you will lose any unsaved data. To end the program now, click End Now.

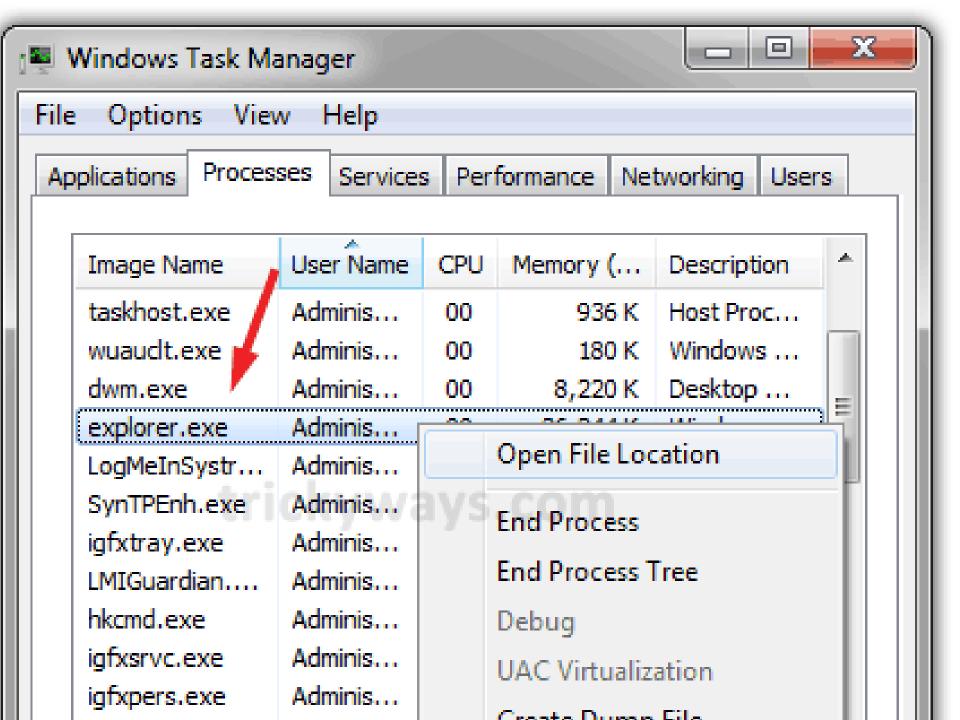
End Now

Cancel







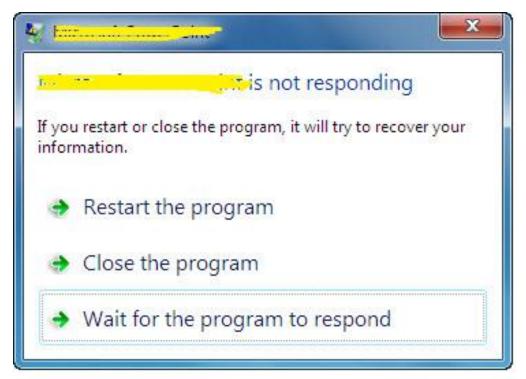


Performance Problems

Widely exist in released software



Mozilla developers fixed 5–60 perf bugs
 monthly over the past 10 years [Jin et al. PLDI 12]



Software Hangs

in daily tools: file managers, office tools, browsers,

...

Software Hangs

- Three major categories
 - Correctness, e.g., infinite loops
 - Blocking operations, e.g., sending files
 - Expensive operations

With **root** of disk (C:\) **selected** in dropdown path selector, attempting to **enable Flat View** under the top-level View menu causes 7-Zip to hang ...

The process had to be forcibly stopped using Windows Task Manager.

contribute to **27%** of 233 hang bugs

[Song et al. DSN 2010]



7-Zip File Manager

Workload-Dependent Performance Bottlenecks (WDPB)

- Expensive operations depending on input workloads, e.g., data processing
- Insight: caused often by workload-dependent loops with expensive operations, e.g.,
 - Temp-object creation/destruction
 - File I/O
 - Ul updates
- Fixed by
 - Spawning new threads
 - Limiting workload/processing sizes

Target Problem: WDPB-Loop Prediction

Traditional **Single**-Profile Setting

Target **Multi**-Profile Setting

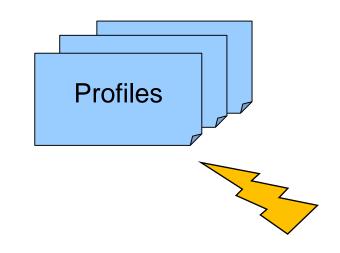
Profile

Test Generation:

Hard to know triggering workload

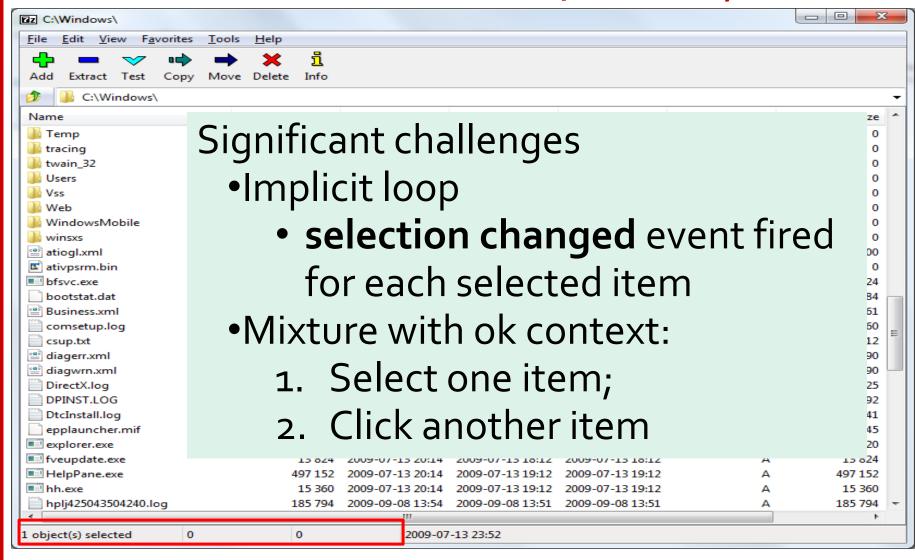
Test Oracle:

Hard to know how slow is slow enough

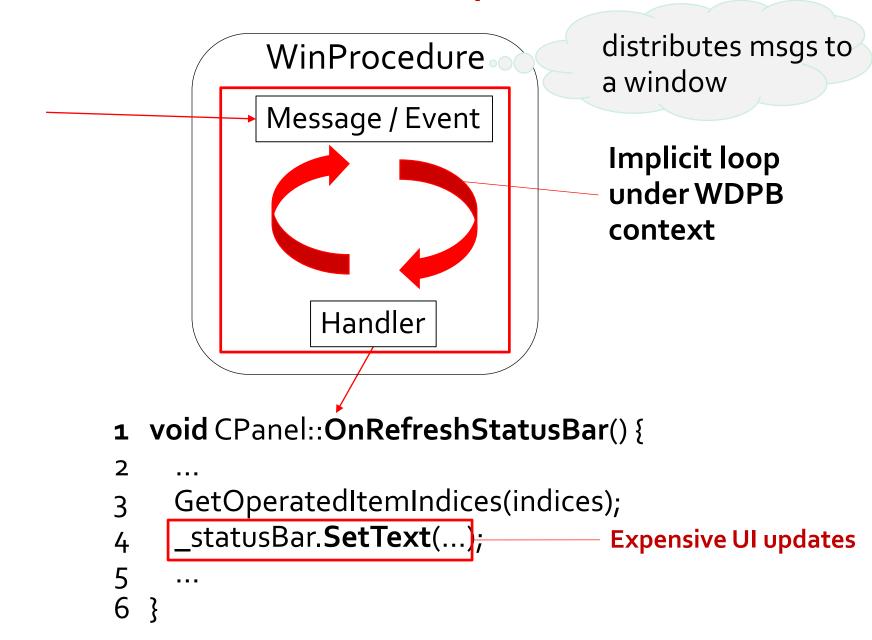


Example WDPB: 7-Zip File Manager

WDPB Context: 1. Select all items; 2. Click any item



Inside the Example WDPB



Background: **StackMine**

[Han et al. ICSE 12]

Perf Debugging in the Large Pattern Matching Bug update Problematic Pattern Bug Database Repository Network **Bug filing** How many issues are still unknown? Key to issue discovery Bottleneck of Trace Storage Trace collection scalability Which trace file should I Trace analysis investigate first?

Proposed Approach: DeltaInfer Context-Sensitive Delta Inference

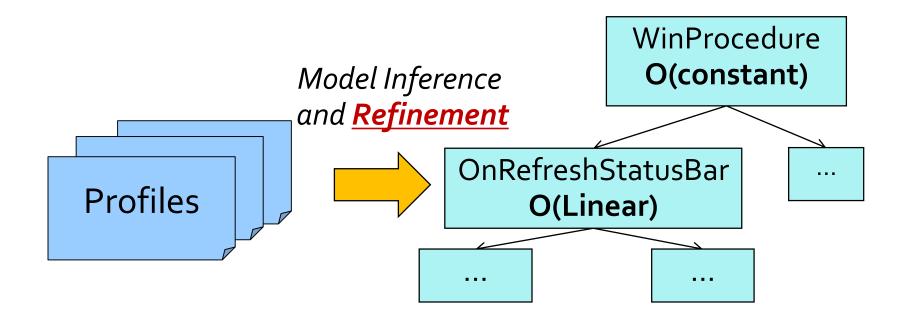
To gain the power of prediction

Temporal Inference: differences between **executions** as complexity models

To identify WDPB loops

Spatial Inference: differences between **program locations** as WDPB loops

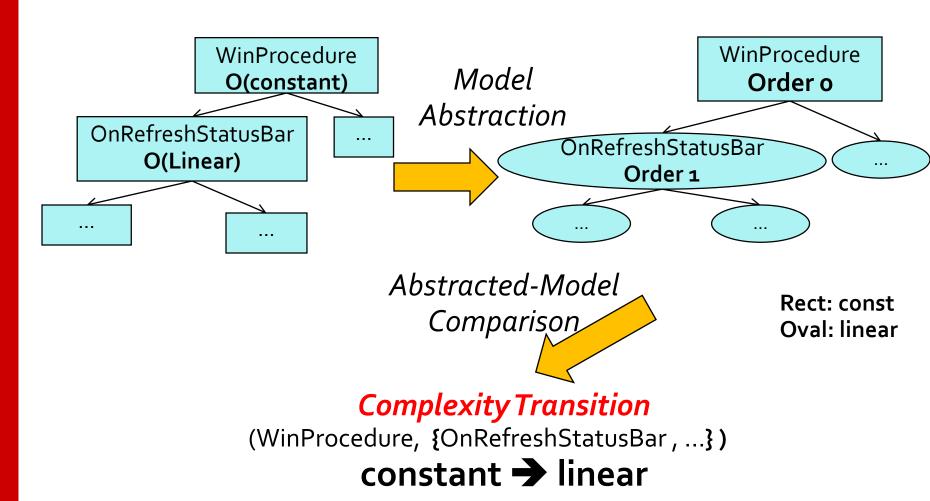
Insight: Temporal Inference



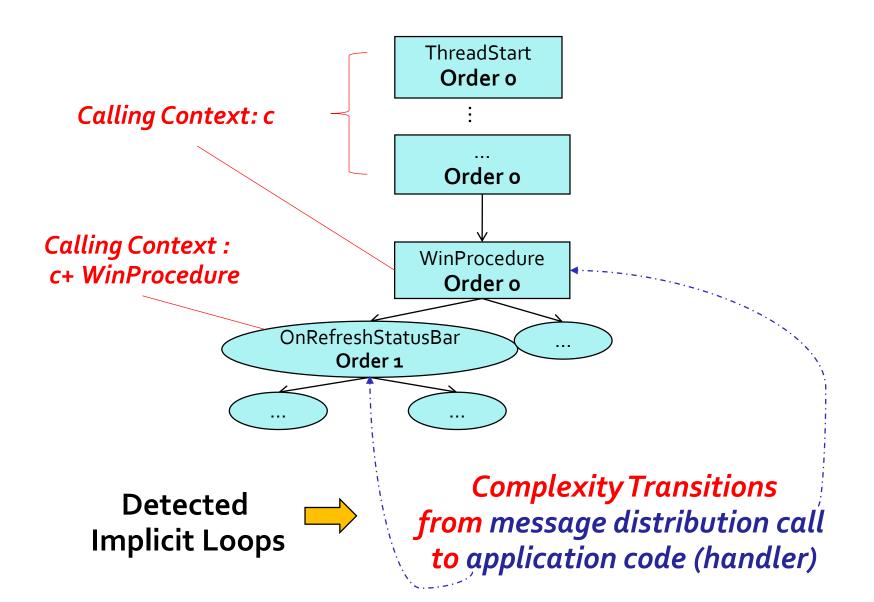
Regression learning + refinement to produce complexity models of program locations

Insight: Spatial Inference

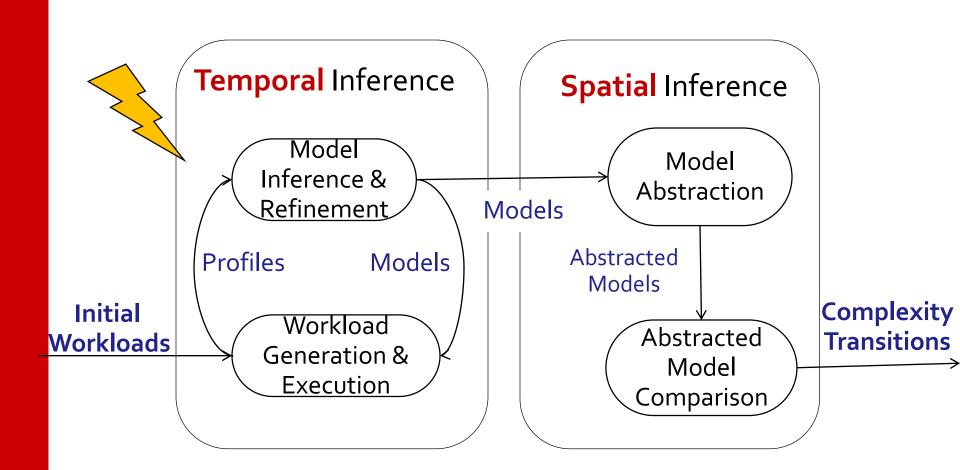
WDPB loop raises complexity models of insideloop locations to higher order



Insight: Context Sensitivity



Overview of DeltaInfer



Workload Generation & Execution

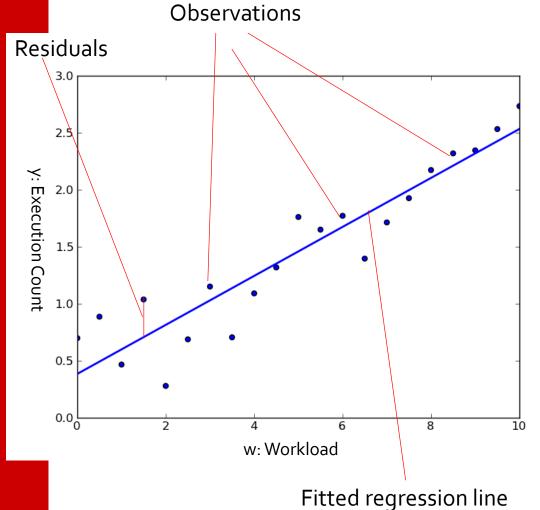
Example scenario: open a file in text editor

- Performance metrics
 - execution time
- Performance-relevant workload parameters
 - # of lines (focused parameter)
 - Rep value range (RVR): [1, 1280]
 - Initial value/variations (sorted/random inputs)
 - # of character

Example workloads

- # of lines (100, 200, ..., 500)
- # of character (100 chars for a line)

Model Inference



- Linear Regression
 - -y=A+Bw
- Power-law Regression

$$-y = Aw^B$$

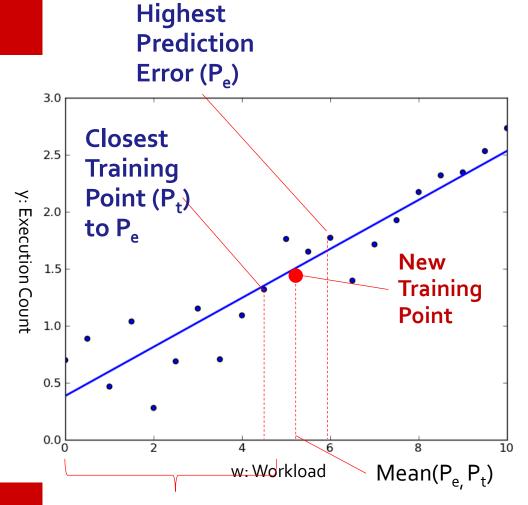
- Quality of the model
 - Correlation coefficient R²

Model Validation

- Model validation measures
 - relative prediction error of inferred models

- Example validation workloads: open a file in text editor
 - Validation value range (VVR): [1,2560]
 - Guideline: >= 2 times larger than the RVR
 - Caveat: too large RVR is not cost-effective

Iterative Refinement



RVR: Representative Value Range

VVR: Validation Value Range

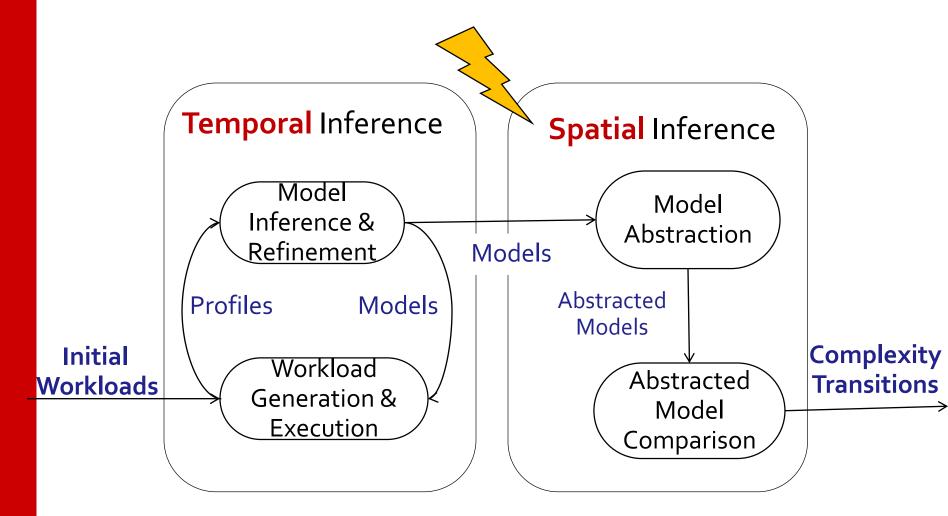
Iterate till

- Accuracy acceptable
- Improvement <threshold

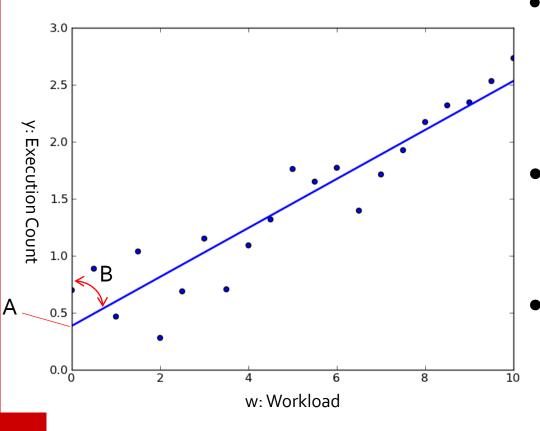
Select a new workload

 Rationale: new workload at highestprediction-error areas improves most

Overview of DeltaInfer

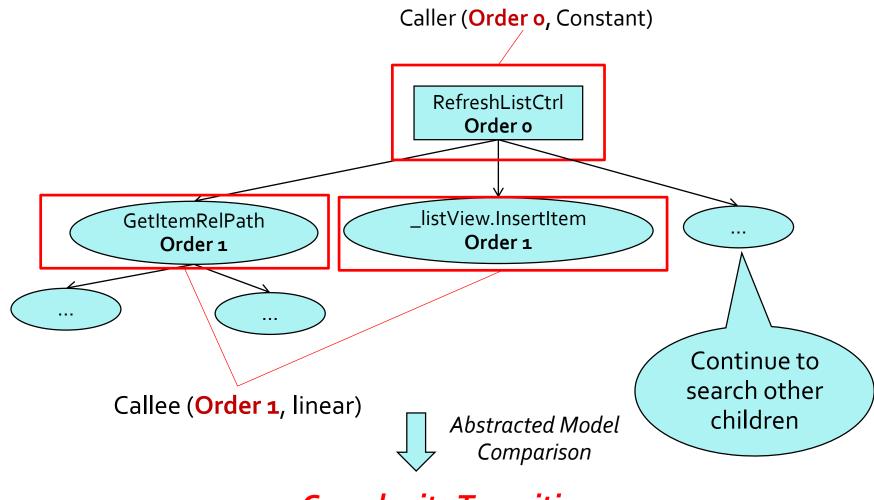


Model Abstraction: Complexity Orders



- Linear model (y = A + Bw)
 - -1, if B > 0
 - o , otherwise
- Power-law model $(y = Aw^B)$
 - Round(B)
- Model w/ R² below threshold_{R₂} (e.g., workload-independent noise)
 - **o**

Inference of Complexity Transitions

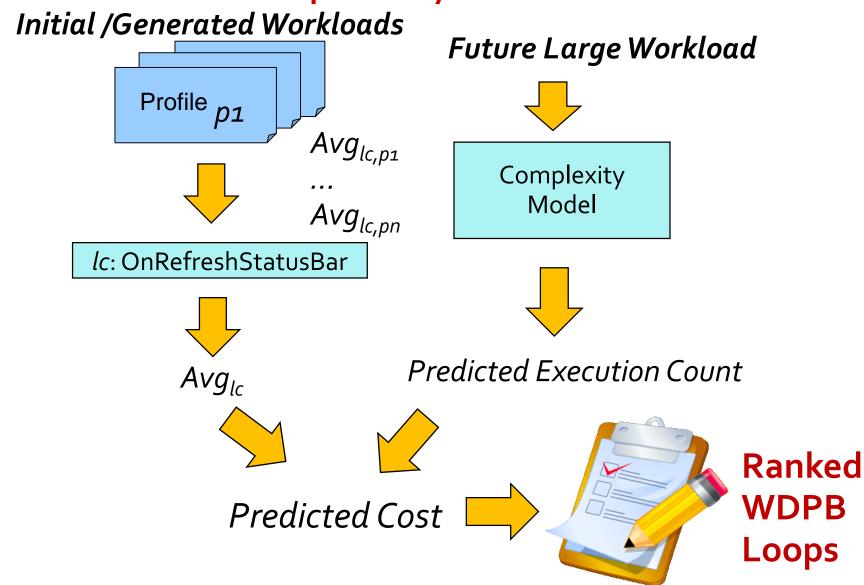


Complexity Transition

(RefreshListCtrl, {GetItemRelPath, _listView.InsertItem, ...})

constant → linear

Cost Prediction of Complexity Transitions



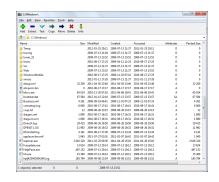
Evaluations of DeltaInfer

 Subjects : open-source GUI applications from SourceForge

• 7-Zip: file manager 7,280 LOC

• Notepad++: text editor 155,300 LOC





- RQs
 - RQ1: Effectiveness of WDPB Identification
 - RQ2: Effectiveness of Model Inference/Refinement
 - RQ3: Effectiveness of Context-Sensitive Analysis

Evaluation Setup - Scenarios

W Daram

lines

	ID	Scenario	w. Param
7-Zip	(S ₁)	Open a folder	# files
	(S ₂)	Rename a file	# files
	(S ₃)	Select all items and then select the first item	# files
	(54)	Create a folder	# files
	(S ₅)	Delete a file	# files
	(S6)	Open a file	# lines
	(S ₇)	Enter a character and save the file	# lines
lotepad++ -	(S8)	Go to the last line	# lines
	(S ₉)	Find a word not present in the file	# chars

Cut and past the first character

No

ID

(S₁₀)

Scanaria

Evaluation Setup – Cont.

- Workload Selection
 - Representative Value Range (RVR)
 - Representative usage [1, 1280]
 - Validation Value Range (VVR)
 - Two times as RVR [1, 2560]
 - Initial workloads
 - Initial workload groups: {20,40,80}, {100,200,400}
- Thresholds
 - Max refinement iterations: 20
 - Threshold for R²: **0.9**; Improvement threshold: **2%**
 - Prediction-error threshold: 5%

RQ1: WDPB Identification

Manually inspect top-rank complexity transitions

 Report identified performance bugs for confirmation from developers

Measure cost coverage of WDPBs as the workloads increase

Example Bugs of 7-Zip

```
Complexity Transition (constant to linear)
(RefreshListCtrl, {GetItemRelPath, _listView.InsertItem, ...}).
HRESULT RefreshListCtrl(...) {
                                                  WDPB loop
  for (UInt32 i = 0; i < numItems; i++) {
                                                Intensive temp-obj
     const UString relPath = GetItemRelPath(i);
                                                creation/destruction
    if (_listView.InsertItem(&item) == -1)
      return E_FAIL;
                                                           Implicit
   listView.SortItems(CompareItems, (LPARAM) this);
                                                           WDPB
                                                           loop!
 Complexity Transition (constant to power-law)
     (listviewProcedure, {CompareItems, ...}).
```

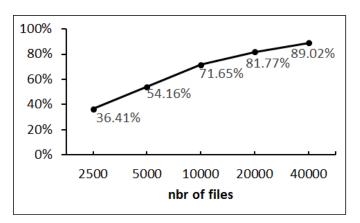
This bug is triggered in S1, S2, S3, S5

Example Bugs of Notepad++

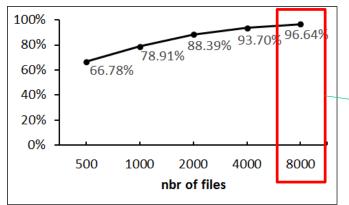
```
WDPB loop
1 bool WrapLines(...) {
    while (lineToWrap < lastLineToWrap) {</pre>
                                                     Expensive
     if (WrapOneLine(surface, lineToWrap)) {
                                                     computation
       wrapOccurred = true;
     lineToWrap++;
10 }
   Complexity Transition (constant to linear)
        (WrapLines, {WrapOneLine, ...}).
```

This bug is triggered in S6, S7 S10

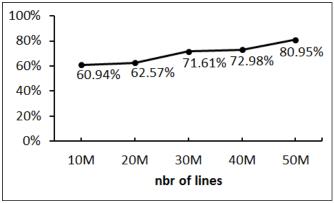
Cost Coverage of WDPBs



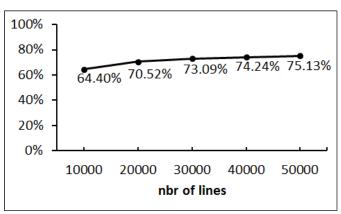
7-Zip FM (S1: open a folder)



7-Zip FM (S3: select all and click any item)



Notepad++ (S6: open a file)



Notepad++ (S9: Find a word)

- Identified WDPBs account for 75%+ of costs → low probability of missing impactful WDPBs
- Cost coverage increases very fast for nested loops

Nested loop!

Bug Confirmations

- 7-Zip (5 bugs)
 - Bugs of **RefreshListCtrl** (S1, S2, S3, S5) confirmed with fix planned for next version
 - Bug of **RefreshStatusBar** (S₄) introduced in release on Aug o₅ and still remaining in latest release on Aug 11.
- Notepad++: (5 bugs)
 - Bug of wraplines (S6) found at the forum
 - Bugs caused by wraplines (S7,S8, S10), pending
 - New bug on search for a not-present word (S9), pending

RQ2: Model Inference and Refinement

Initial Workloads

ID	# Ite.	# WorkL	E. I.	E. E.
(S1)	4	4	35.95	0.62
(S2)	4	4	62.31	0.47
(S3)	4	4	29.68	0.85
(S4)	4	4	4.71	0.20
(S5)	3	3	5.94	0.18
(S6)	6	6	536.74	8.62
(S7)	5	5	455.00	7.69
(S8)	7	7	17.12	5.51
(S9)	4	4	138.36	1.83
(S10)	7	7	7.38	1.86

After Refinement

Avg relative errors of inferred complexity models

- 5 iterations (7 workloads) to reach avg relative err 2.8%
- Insensitive to potential variations of initial workloads

RQ2: Prediction Error of Cost

ID	10 (%)	20 (%)	50 (%)
(S ₁)	3.18	4.45	6.16
(S ₂)	2.98	4.07	5.55
(S ₃)	*1.40	*1.60	*1.86
(S ₄)	1.65	2.29	3.08
(S ₅)	1.58	2.19	2.95
(Ave(7-Zip))	*2.35	*3.25	*4.44
(S6)	18.51	6	47.24
(S ₇)	16.84	5	36.28
(S8)	16.80	7	35.23
(S ₉)	11.15	4	39.09
(S10)	10.79	7	24.63
(Ave(Notepad++))	14.82	20.97	36.49

X RVR Upper Bound

Developers optimize message processing during idle time

Prediction error

- 4.4% (7-Zip file manager): excluding S3
- 36.5% (Notepad++): robust even under complex situations

RQ3: Context Sensitivity

ID	DeltaInfer		# L.		InSen. # Missed by InSen			
(S ₁)	11		10		521		6	Real WDPB loops
(S ₂)	21		19		579	_	12	- Real World Bloops
(S ₃)	17		16		486		10	
(54)	21		19		640		12	
(S ₅)	22		20		546		12	
(S6)	10		10		509		3	
(S ₇)	29		14		877		6	
(S8)	10		10		526		5	
(S ₉)	20		20		131		0	
(S10)	12		11		861		3	

- Context helps reduce false positives & negatives
 - No context: > 90% of identified WDPB loops being false positives
 - No context: 40% of DeltaInfer-identified WDPB loops being missed
- Context helps achieve only 14% of identified WDPB loops being false positives (top/low-level sys lib calls)

Conclusion

- Predictive approach for WDPBs: context-sensitive delta inference
 - Temporal inference complexity models
 - **Deltas** of different **executions** (workloads)
 - Spatial inference → complexity transitions
 - Order deltas of different locations

Evaluations: effectively identifies impactful
 WDPBs (for causing 10 performance bugs)



Microsoft Research

Thank You!



Supported in part by NSF grants CCF-0845272, CCF-0915400, CNS-0958235, CNS-1160603

Conclusion

- Predictive approach for WDPBs: context-sensitive delta inference
 - Temporal inference complexity models
 - **Deltas** of different **executions** (workloads)
 - Spatial inference → complexity transitions
 - Order deltas of different locations

Evaluations: effectively identifies impactful
 WDPBs (for causing 10 performance bugs)

Discussion

 Generalization to Other Types of Applications

Multiple Workload Parameters

Value-Dependent Performance Bottlenecks

Scalability of Scenario-Based Profiling

Software Performance

- An important quality of software
 - A kind of non-functional requirement
 - Characterized by the amount of work accomplished given time and resources
- Software performance matters

- Software functionality and size grows faster than



been done elated mista

avoid

Tablet

http://betanews.com/2012/05/09/software-performData Centers/

Two Significant Challenges in GUI Applications

- Complex contexts
 - GUI applications are event-driven applications
 - A program location may exhibit different complexities under different contexts

- Implicit WDPB loops
 - Event handlers can be invoked repetitively
 - E.g., selection change events for all items
 - No explicit loop statements
 - Cause challenges to manual inspection and static analysis

Limitations of Traditional Approaches

- Traditional approaches
 - Performance testing (blackbox-random testing or manual testing)
 - Profiling (call-tree profiling and callstack sampling)

41 out of 109 studied performance bugs are due to wrong assumption of workloads. [Jin et al. PLDI 2012]

- Two major issues
 - Insufficiency
 - WDPBs may not surface on given workloads
 - Workload specifications are usually missing or outdated
 - Incompleteness:
 - WDPBs may overshadow other WDPBs

Least-Squares Regression

- Linear Regression
 - Infers: y = A + Bw
 - Minimizes: $Q(A,B) = \sum_{i=1}^{k} (y_{l_c,i} (A + Bw_i))^2$
- Power-law Regression
 - Infers: $y = Aw^B$
 - Minimizes: $Q(A, B) = \sum_{i=1}^{k} (y_{l_c, i} (Aw_i^B))^2$
- How good does the model fit the data points?
 - Correlation coefficient

•
$$R^2 = \frac{(\sum_{i=1}^k wy - k\overline{wy})^2}{(\sum_{i=1}^k w^2 - k\overline{w}^2)(\sum_{i=1}^k y^2 - k\overline{y}^2)}$$
,

• \overline{w} is the mean of w, \overline{y} is the mean of y

A Few Definitions

- Application A, location l and cost y
- Call graph G(E,V), calling context c, and execution profile P
- k-profile Graph: an annotated call graph, G(E, V), where a location I with its corresponding vertex is annotated with a vector of counters for I on k workloads for each of its calling context c.

Complexity Transitions

- A pair (n,M), such that:
 - 1. n is a vertex (method) in the k-profile graph and M
 is a subset of children vertices (callees) of n;
 - 2. $f_{n,c}(W)$ is the complexity model of n under the calling context c, and $f_{li,ci}(W)$ is the complexity model of the location l_i , where l_i is a location in M and the calling context c_i is c concatenated with n.
 - $-3. O(f_{li}, c_i(W))$ is at least 1 more than $O(f_{n,c}(W))$;
 - 4. $\forall l_i, l_j \in M$, $i \neq j$, $O(f_{l_i,c_i}(W)) = O(f_{l_j,c_j}(W))$.

Model Inference and Refinement

```
1: pc = -1 // previous model count
 2: e_{pre} = -1 // previous error
                                 = ite + 1 do
      kG = AlignProfiles(P)
      r - GetWorkloads(P)
     M = RegressionLearning(x, kG)
7:
      for all vp in VP do
 8:
9:
         w = GetWorkload(vp)
10:
         e_{vp} = 0.0
         for all m in M do
11:
12:
           r_w = Predict(w, m)
           a_m = GetActual(vp, m)
13:
           e_{vp} = e_{vp} + \frac{Abs(a_m - r_w)}{a_m}
14:
15:
         e_M = e_M.Add(\frac{e_{vp}}{M.Count})
16:
17:
      end for
      e_{total} = \frac{Sum(e_M)}{VR.Count}
18:
      if e_{pre} - e_{total} < threshold_{imp} then
19:
         break // improvement is below the threshold
20:
21:
      end if
22:
      if e_{total} < threshold_e AND pc == M.Count then
23:
         break // accuracy is acceptable
24:
      _{
m else}
25:
        np = NewWorkload(P, VP, e_M)
26:
        P = P.Add(np)
27:
         pc = M.Count
         e_{pre} = e_{total}
29:
      end if
30: end for
31: return M
```

Align Profiles

- Align locations using calling contexts
- Extract execution vector for each location under each calling context

Regression Learning

Model Validation

Termination Checks

Select New Workloads

 Assumption: a new workload at the area with the highest prediction error improves most

Cost Prediction of Complexity Transitions

- Compute $avg_{lc,p}$ for each location l_c on each profile p
 - E.g., for p_1 , $Cost(refreshList_c) = 1s$, $ExeCount(refreshList_c) = 100$, $avg_{lc,p} = 1/100$ s
- Compute $avg_{lc} = average(avg_{lc,p})$
- Given a workload value w, $pred_{lc,w} = f_{lc}(w)$
- Get $Cost_{lc,w} = pred_{lc,w} * avg_{lc}$