



StackMine – Performance Debugging in the Large via Mining Millions of Stack Traces

Shi Han¹, Yingnong Dang¹, Song Ge¹, Dongmei Zhang¹, and Tao Xie²



Software Analytics Group, Microsoft Research Asia¹

North Carolina State University²

June 6th, 2012

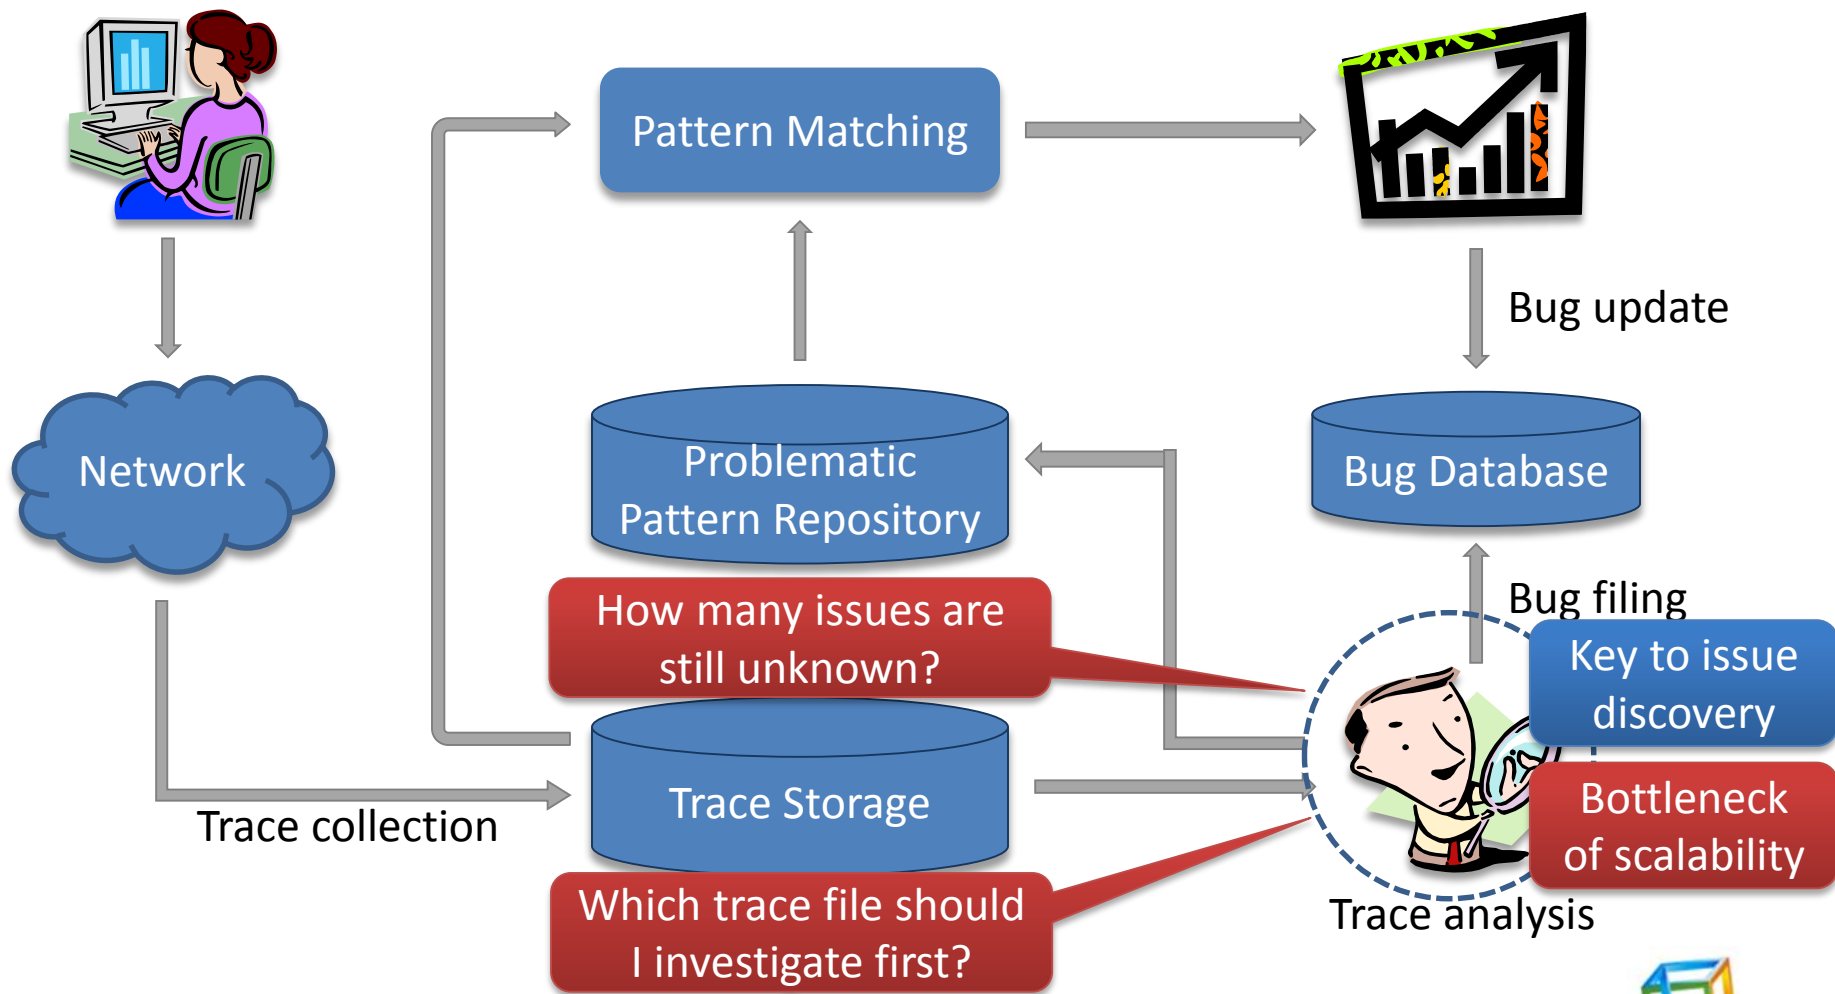
Performance Issues in the Real World

- One of top user complaints
- Impacting large number of users every day
- High impact on usability and productivity



Given **limited** time and resource **before** software release, **development-site** testing and debugging become **insufficient** to ensure satisfactory software performance.

Performance Debugging in the Large





Problem Definition

Input: Runtime traces

collected from millions of users

Output: Program execution patterns
causing the most impactful performance
problems

inspected by performance analysts






Goal

Conduct **systematic discovery & analysis of program execution patterns**

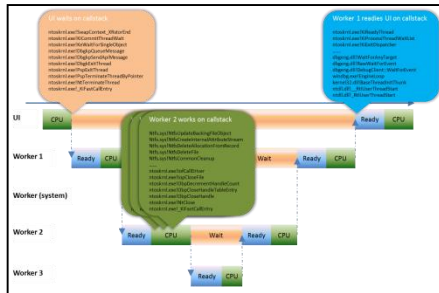
- Efficient handling of large-scale trace sets
- Automatic discovery of new patterns
- Effective prioritization of investigation

Challenges



Large-scale trace data

- TBs of trace files and increasing
- Millions of events in single trace stream



Highly complex analysis

- Numerous program runtime combinations triggering performance problems
- Multi-layer runtime components from application to kernel being intertwined



Combination of expertise

- Generic machine learning tools without domain knowledge guidance do not work well

What
brow

Wait callstack

```
ntdll!UserThreadStart
Browser! Main
...
Browser!OnBrowserCreatedAsyncCallback
...
BrowserUtil!ProxyMaster::GetOrCreateSlave
BrowserUtil!ProxyMaster::ConnectToObject
...
rpc!ProxySendReceive
...
wow64!RunCpuSimulation
wow64cpu!WaitForMultipleObjects32
wow64cpu!CpupSyscallStub
...
```

Wait callstack

```
ntdll!UserThreadStart
Browser! Main
...
ntdll!LdrLoadDll
...
nt!AccessFault
nt!PageFault
...
```

ReadyThread callstack

```
nt!KiRetireDpcList
nt!ExecuteAllDpcs
...
nt!IoPcCompleteRequest
...
nt!SetEvent
...
```

UI thread



Underlying Disk I/O

Unexpected long execution

Worker thread




CPU sampled callstack

```
ntdll!UserThreadStart
...
ntdll!WorkerThread
ole!CoCreateInstance
...
ole!OutSerializer::UnmarshalAtIndex
ole!CoUnmarshalInterface
...
```

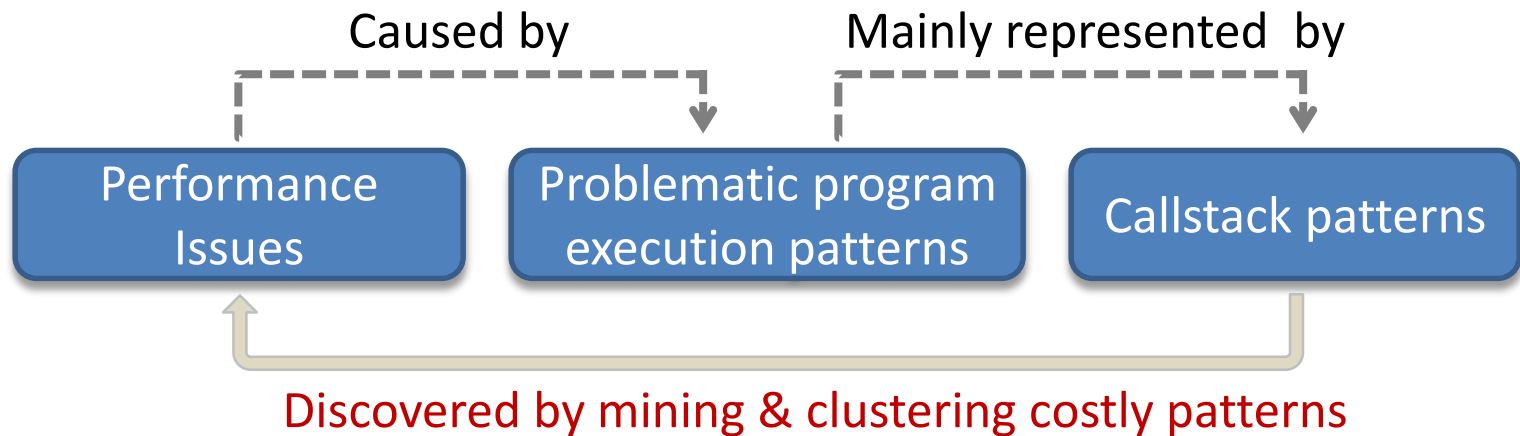
ReadyThread callstack

```
ntdll!UserThreadStart
...
rpc!LrpcIoComplete
...
user32!PostMessage
...
win32k!SetWakeBit
nt!SetEvent
...
```




Approach

Formulated as a callstack mining and clustering problem




Approach – Workflow

Trace streams



AOI Extraction

Callstacks



Ranked cluster

Cluster	Hits
1	94,279
2	51,107
3	35,536
...	...

Pattern Clustering

Sequence Mining

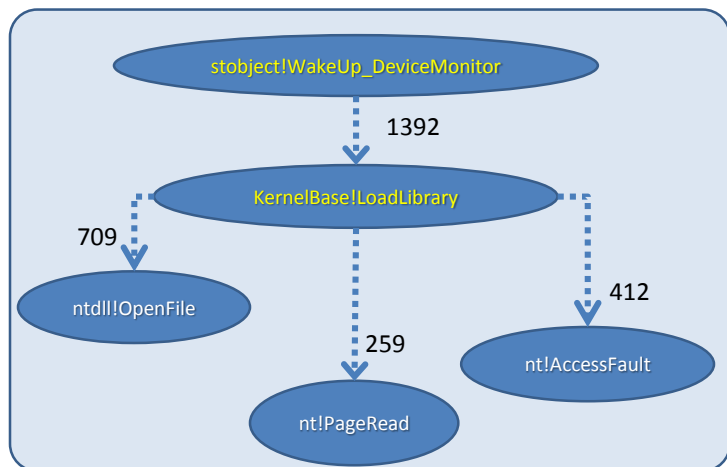
Ranking

Ranking

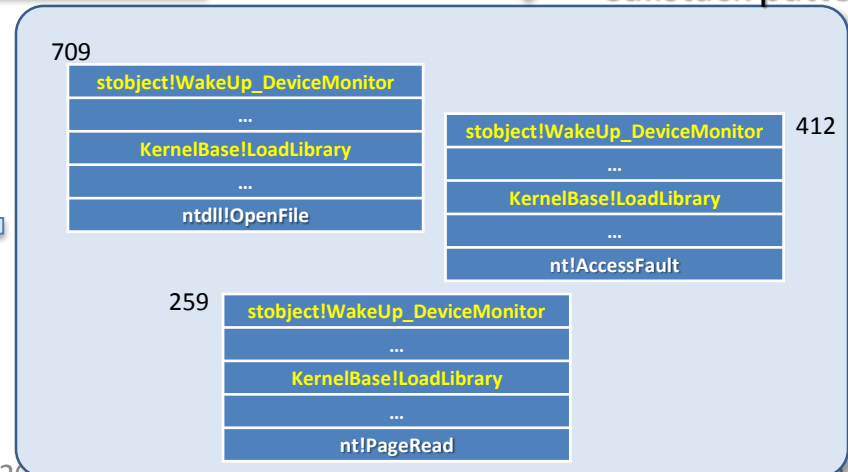
Ranking

Callstack patterns

Pattern clusters




Pattern Clustering



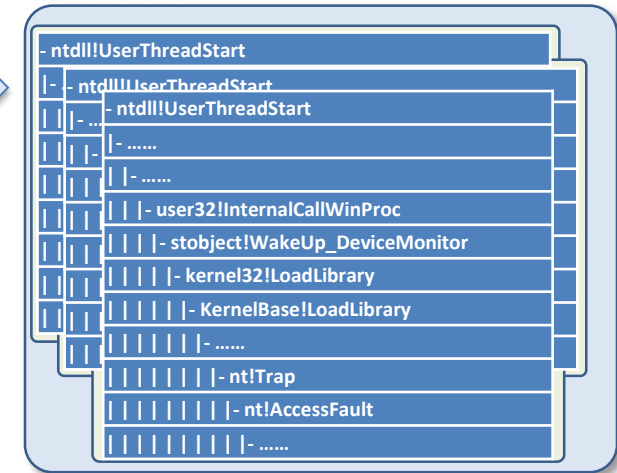
Approach – AOI Extraction

Trace streams



AOI Extraction

Callstacks

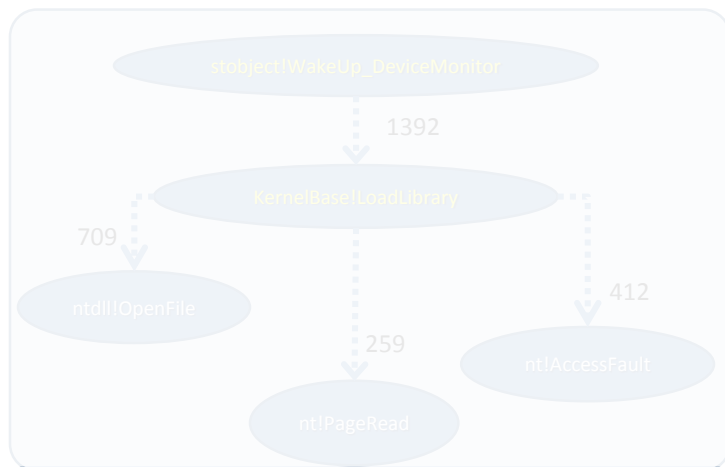


Ranked cluster

Cluster	Hits	Total Wait time (ms)
1	94,279	927,824
2	51,107	561,416
3	35,536	3,051,307
...

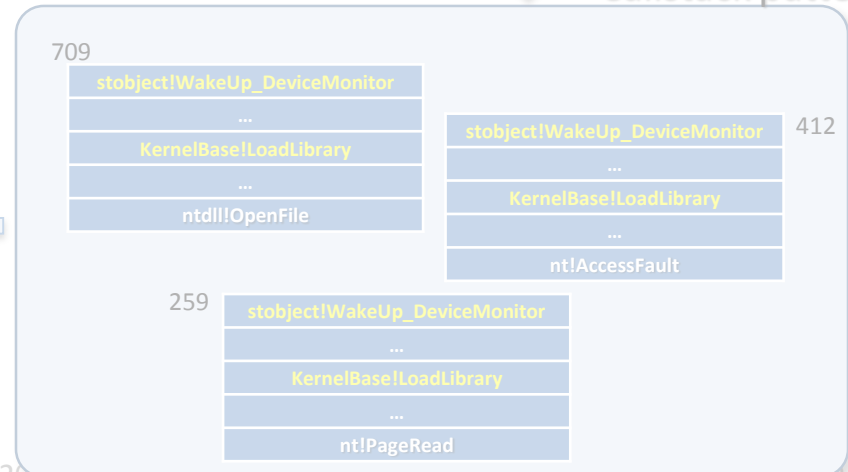
Ranking

Pattern clusters



Sequence Pattern Mining

Callstack patterns



Pattern Clustering



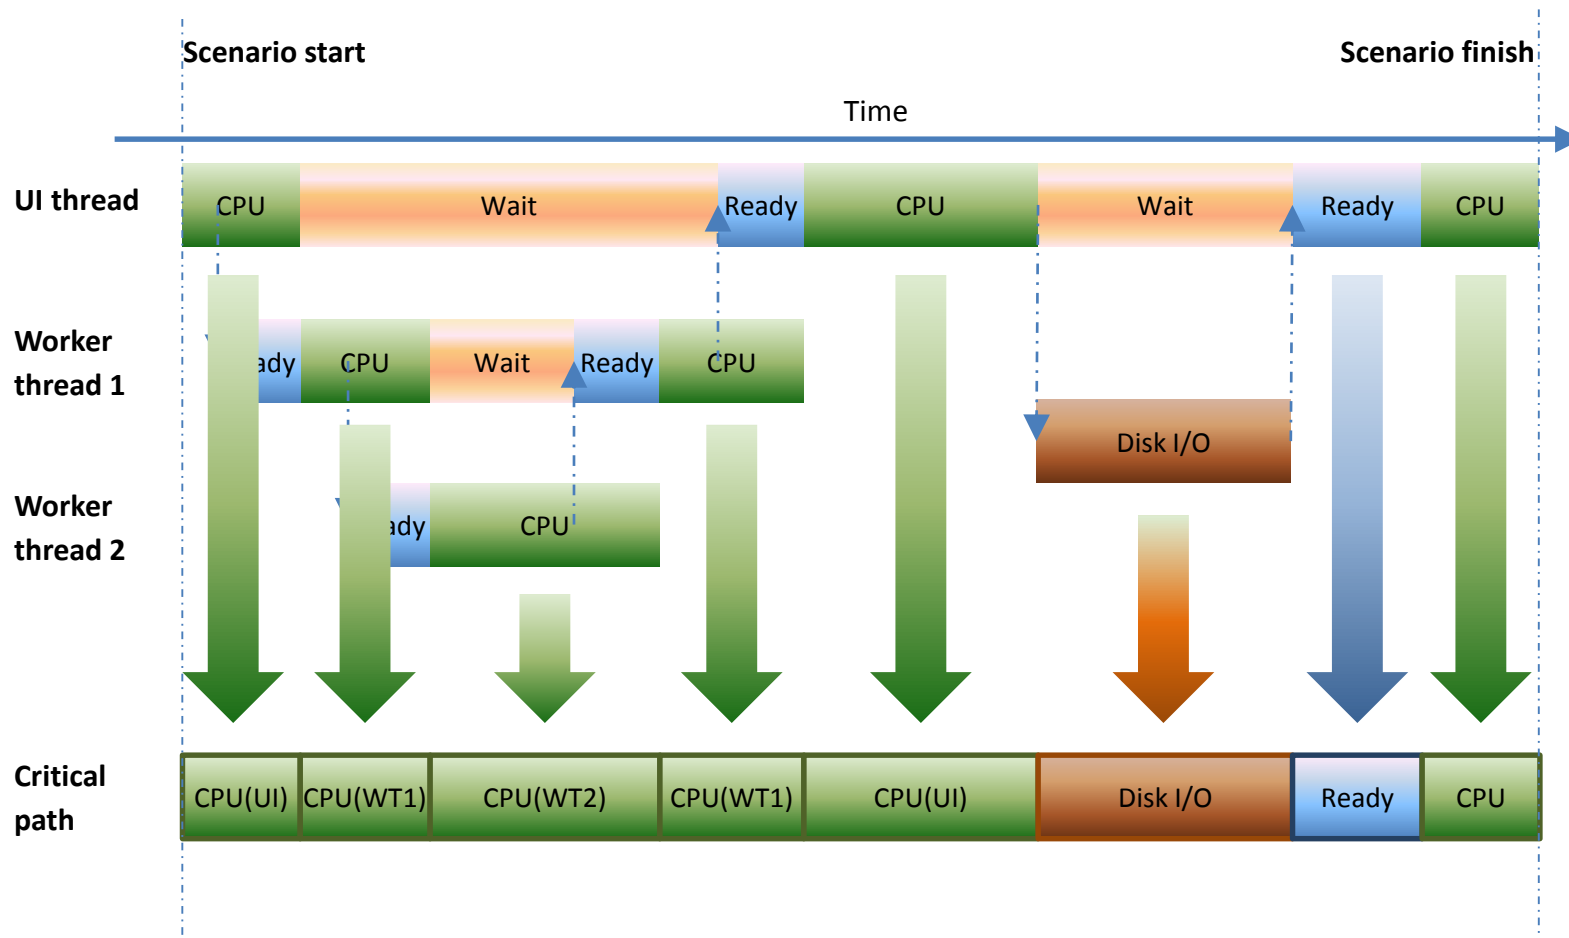
Approach – AOI Extraction

Motivation

Runtime traces capture both

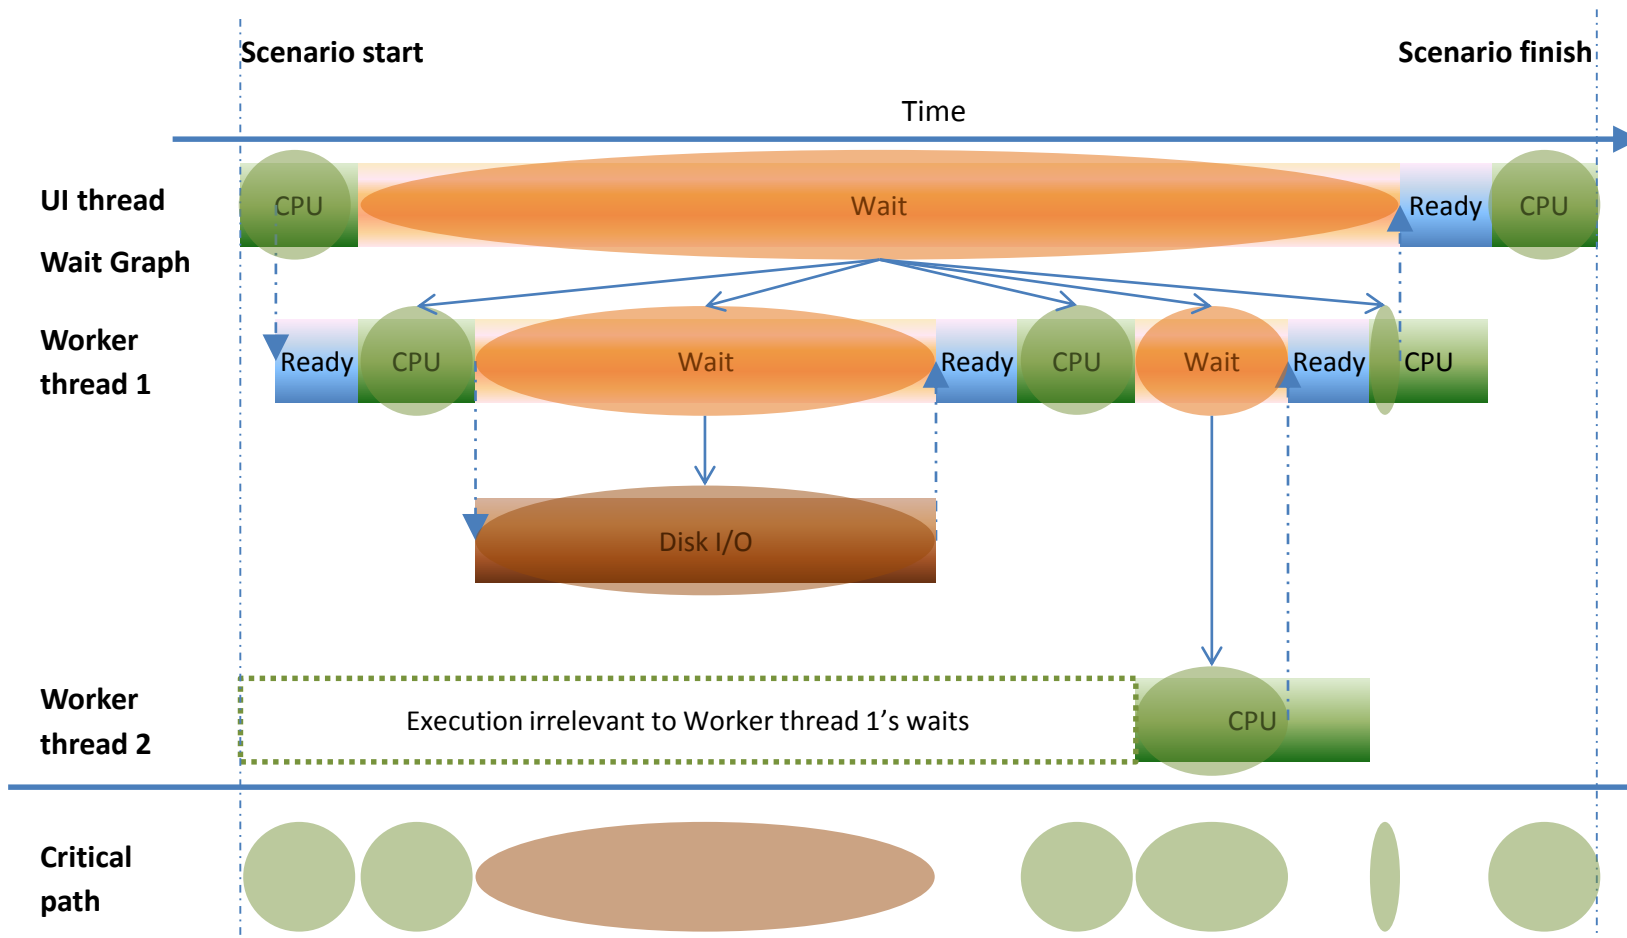
- Relevant executions for performance issue
 - E.g., executions relevant to browser-tab creation
- Irrelevant executions for performance issues
 - E.g., executions of concurrently executed IM
 - Noisy data for mining
 - Huge investigation scope induced

Approach – AOI Extraction Critical Path




Approach – AOI Extraction

Wait Graph



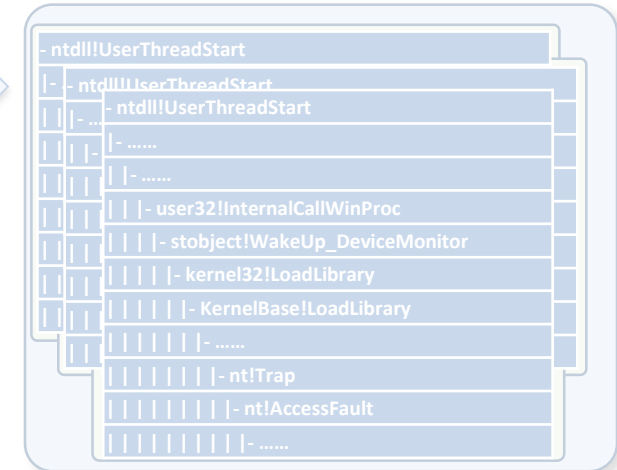
Approach – Callstack Pattern Mining

Trace streams



AOI Extraction

Callstacks

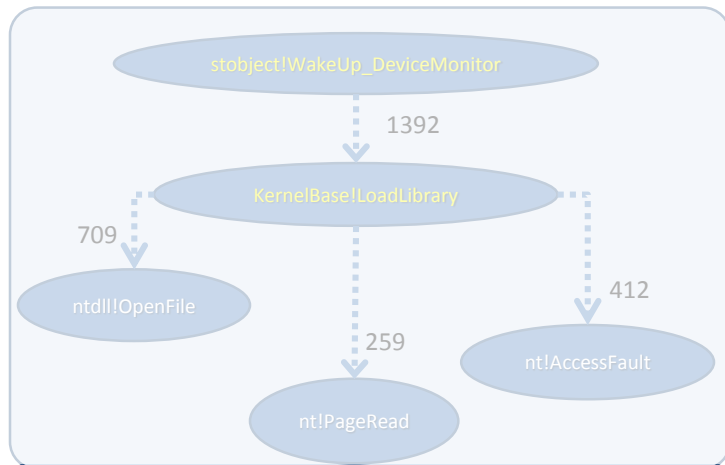


Ranked cluster

Cluster	Hits	Total Wait time (ms)
1	94,279	927,824
2	51,107	561,416
3	35,536	3,051,307
...

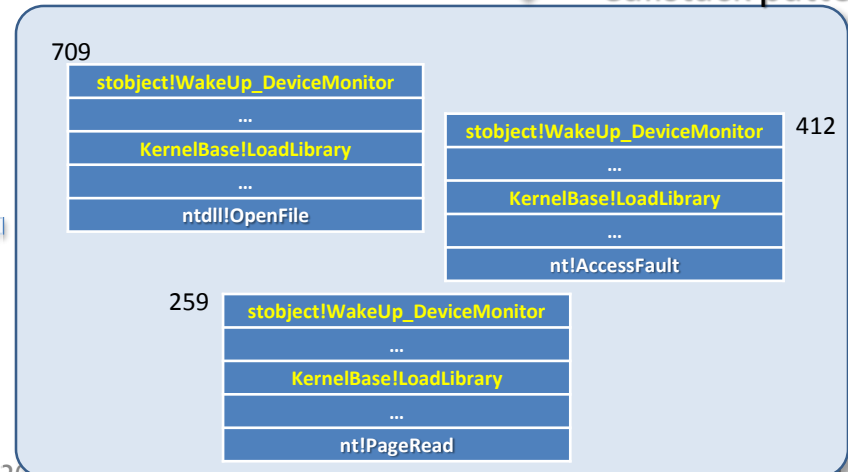
Ranking

Pattern clusters



Sequence Pattern Mining

Callstack patterns



Pattern Clustering

Approach – Callstack Pattern Mining

Costly Sequence Pattern

- Example

Sequence	<u>Wait time</u>
A B	4 ms
A B C D F	1 ms
A B C E F	3 ms
A B G F	2 ms

Cost threshold T

5 ms (or 50% of total)


Costly sequence pattern miner

All Patterns	Support	Closed Patterns	Support	Maximal Patterns	Support
A, B, AB	10	A, B , AB	10	A, B , AB	10
F, AF, BF, ABF	6	F, AF, BF , ABF	6	F, AF, BF , ABF	6

- Non-consecutive costly sub-sequence as callstack pattern
- Costly maximal patterns are compact

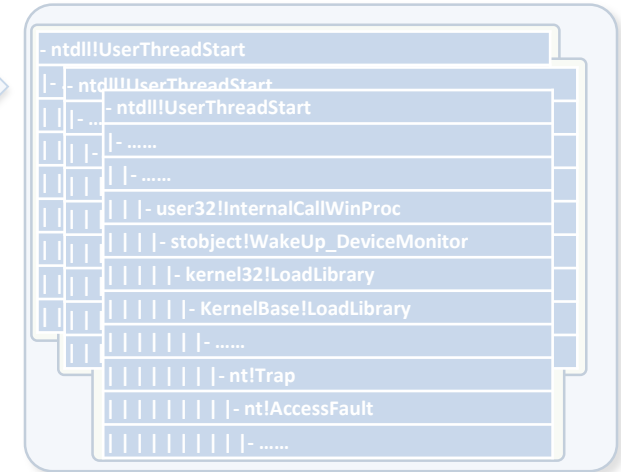
Approach – Callstack Pattern Clustering

Trace streams



AOI Extraction

Callstacks




Ranked cluster

Cluster	Hits	Total Wait time (ms)
1	94,279	927,824
2	51,107	561,416
3	35,536	3,051,307
...

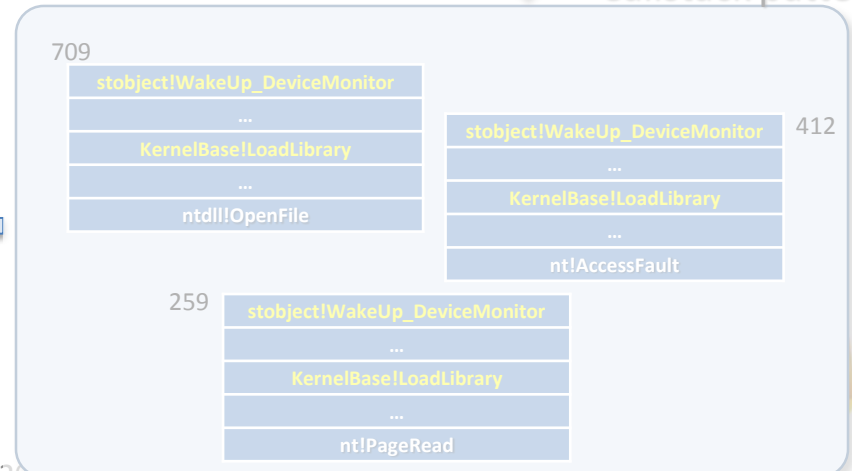
Ranking

Pattern clusters



Sequence Pattern Mining

Callstack patterns



Pattern Clustering

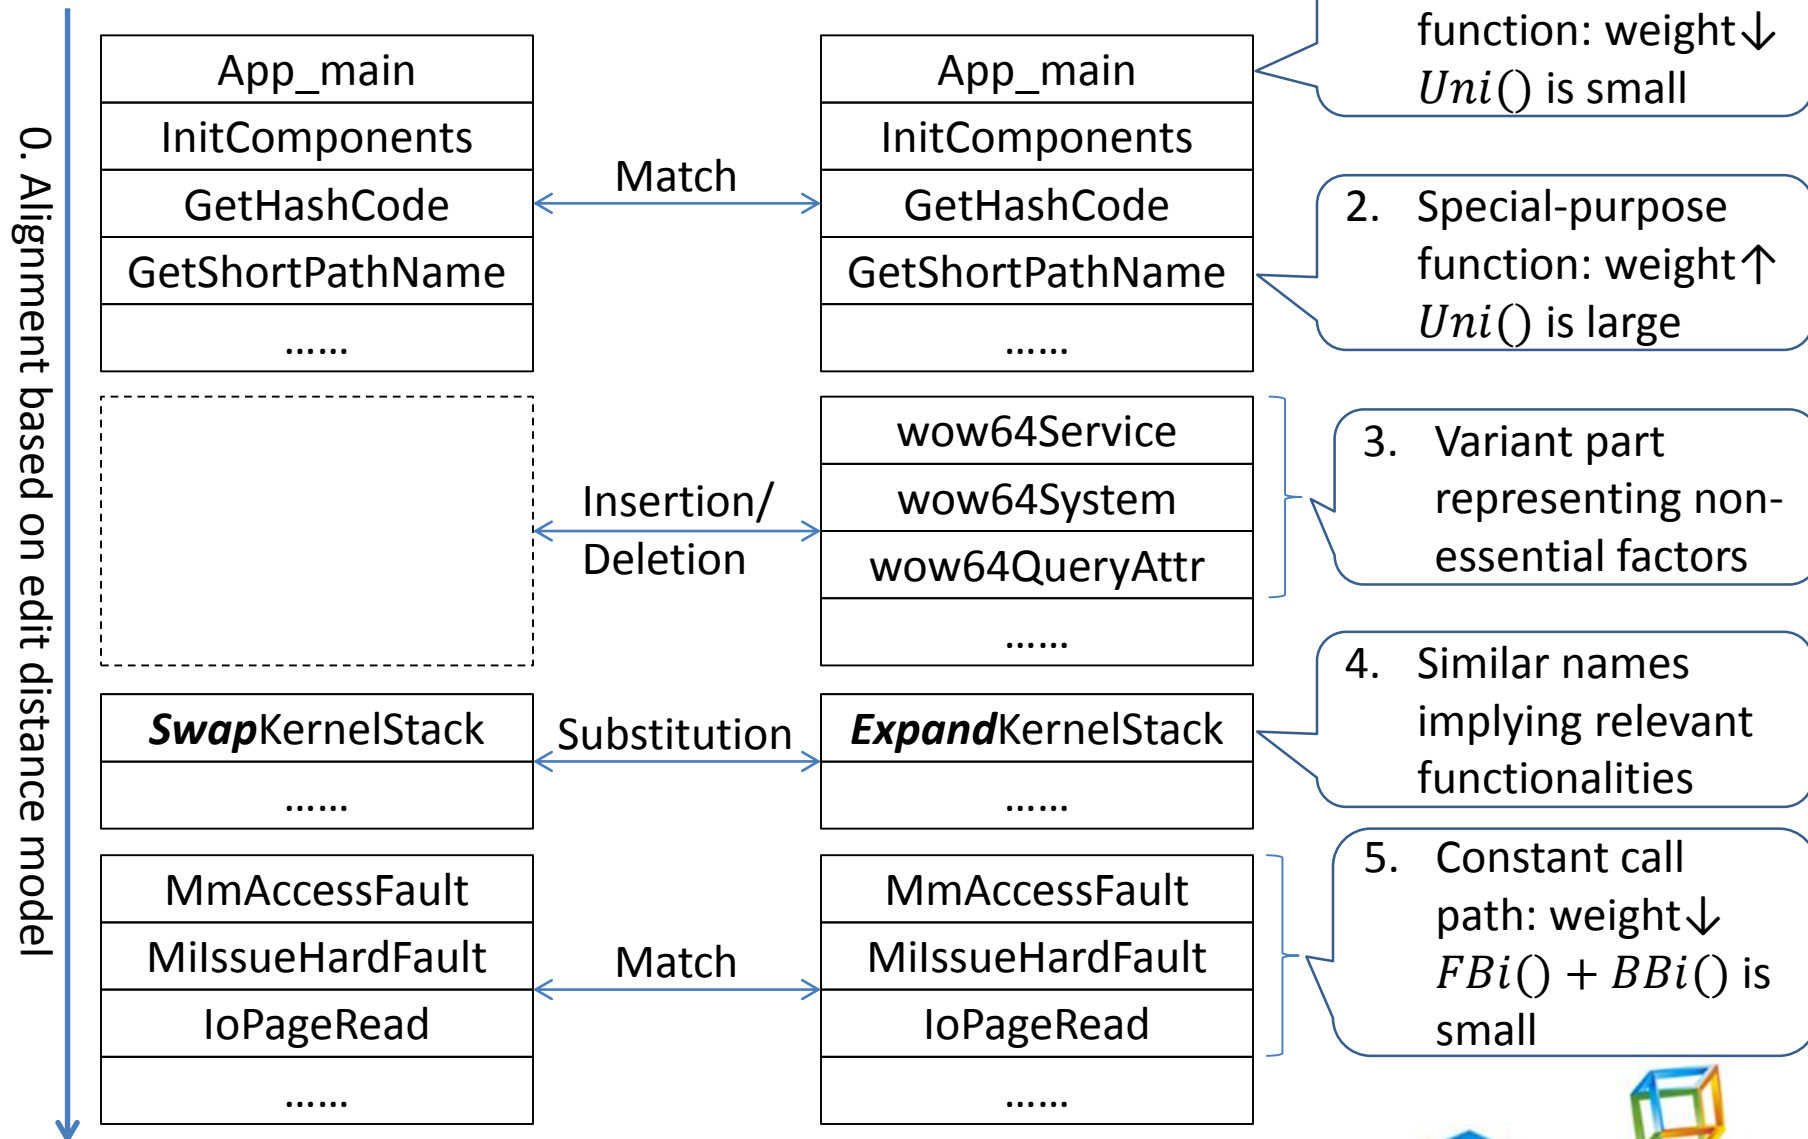


Approach – Callstack Pattern Clustering

- Motivation
 - Same issue often reflected by variant patterns
 - Defect often hidden in invariant parts of variant patterns
- Goal
 - Precise measurement of issue impact for better prioritization
 - Comprehensive issue representation with pattern variations for quick and precise fixing

Approach – Callstack Pattern Clustering

Similarity Model





Technical Highlights

- Machine learning for system domain
 - Formulate the discovery of problematic execution patterns as callstack mining & clustering
 - Systematic mechanism to incorporate domain knowledge
- Interactive performance analysis system
 - Parallel mining infrastructure based on HPC + MPI
 - Visualization aided interactive exploration





Evaluation – Windows 7 Study

- Task: since Dec 2010, a continued effort to improve Windows performance
 - **Analysts** from one performance analysis team for Microsoft Windows
 - **Hunt** for hidden performance bugs that caused common impact on Windows Explorer UI response
 - **Based** on over 6,000 trace streams
- Data
 - 921 qualified out of 1,000 randomly sampled trace streams
 - 181 million callstacks in total
 - 140 million wait callstacks
 - 41 million CPU sampled callstacks





Evaluation – Windows 7 Study

Research Questions

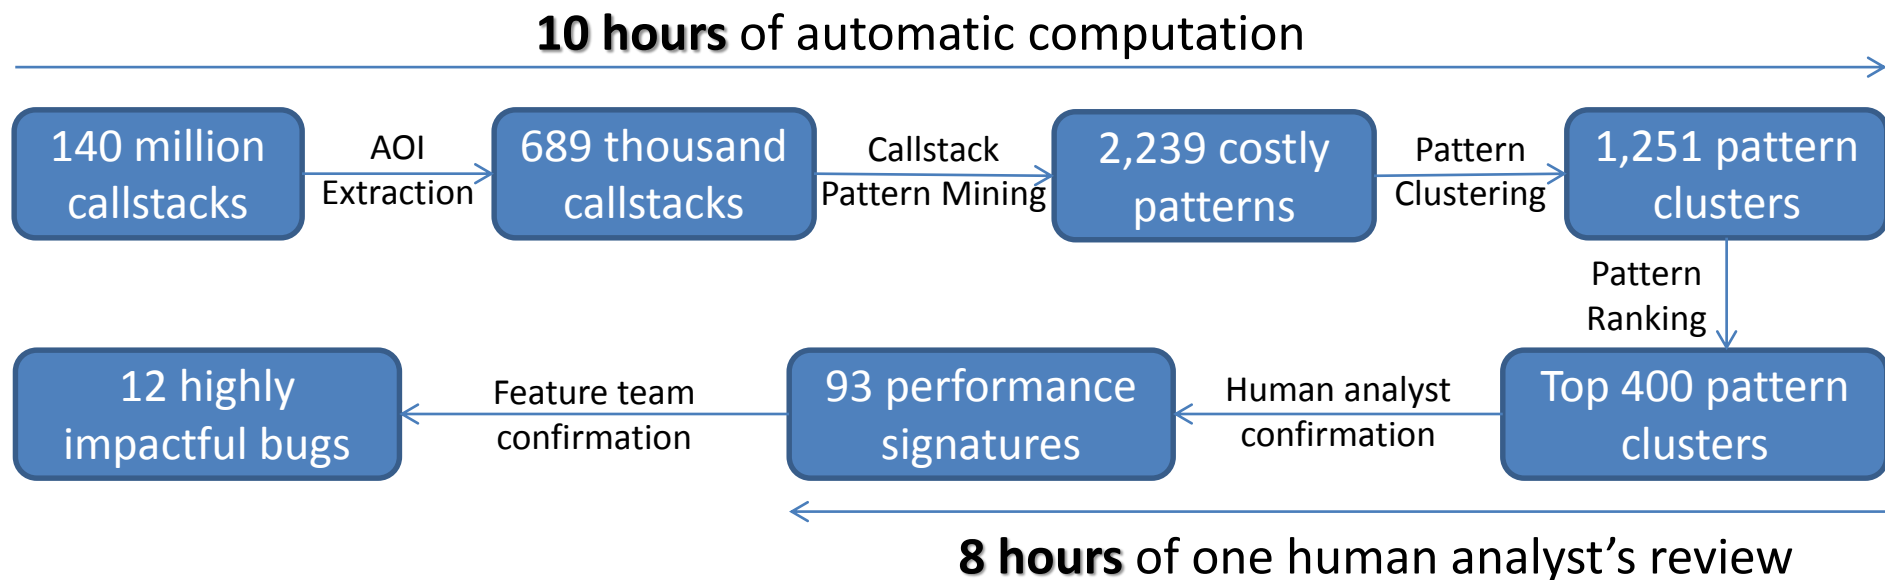
- *RQ1.* How **much** does StackMine **improve** practices of performance debugging in the large?
- *RQ2.* How **well** do the derived performance signatures **capture** performance bottlenecks?
-
- *RQ3.* How **much** does StackMine **outperform** alternative techniques?



Evaluation – Windows 7 Study

RQ1. Overall Improvement of Practices

- Traditional approach – would take 20~60 days
- Using StackMine – 18 hours





Evaluation – Windows 7 Study

RQ2. Performance Bottleneck Coverage

- Performance Bottleneck Coverage (PBC) of a set of performance signatures

$$PBC = \frac{\text{Total time of performance signatures}}{\text{Total time of collected trace streams}}$$

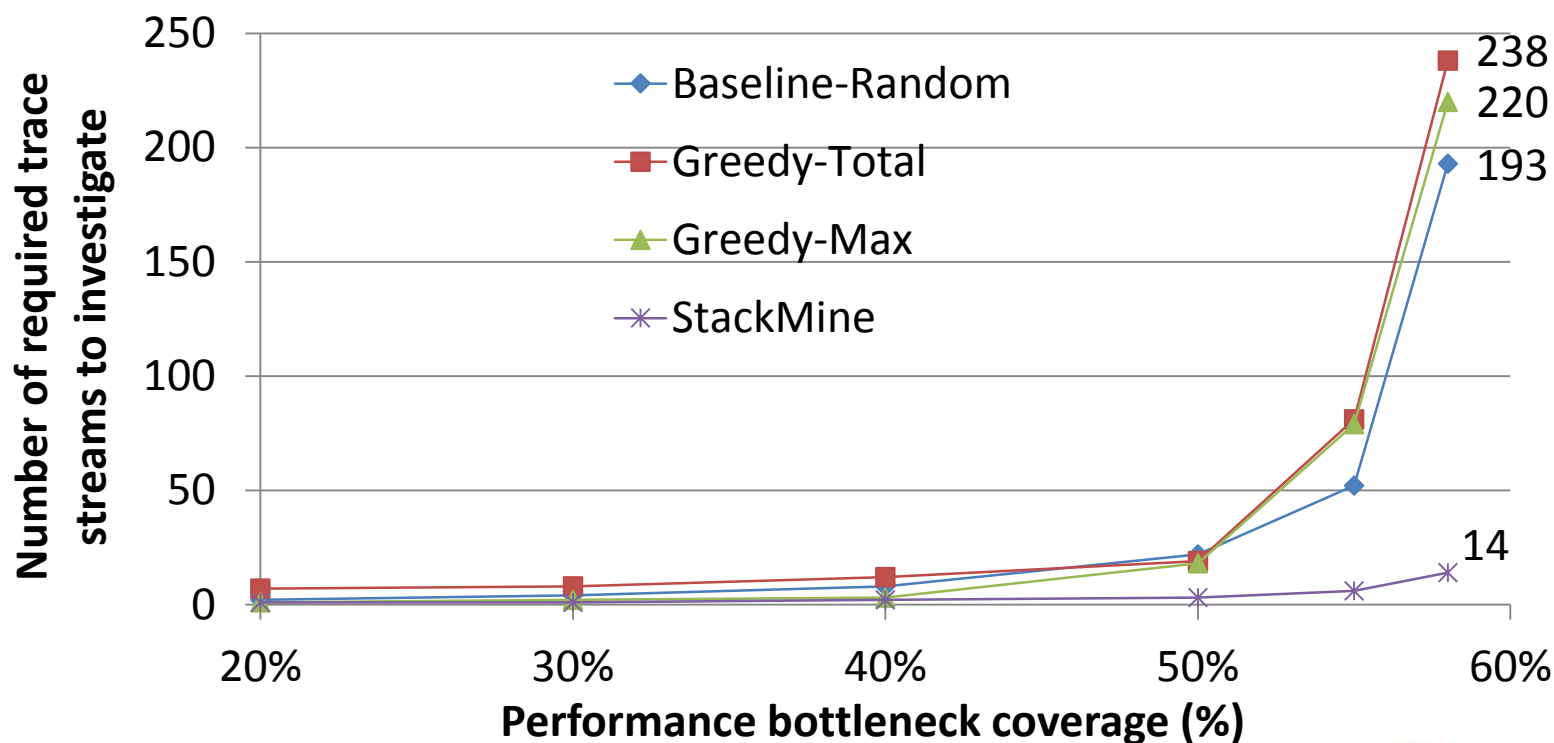
- The higher PBC achieved, the lower possibility that high-impact performance bugs remain not captured
- **58.26%** PBC achieved by the 93 signatures



Evaluation – Windows 7 Study

RQ3. Comparison with Alternative Techniques

StackMine requires only **7.2%**, **5.8%**, and **6.3%** of trace streams required by the other three techniques



Impact



*“We believe that the MSRA tool is highly valuable and much more efficient for mass trace (100+ traces) analysis. For **1000 traces**, we believe the tool **saves us 4-6 weeks** of time to create new signatures, which is quite a significant productivity boost.”*

- from Development Manager in Windows

Highly effective new issue discovery on
Windows mini-hang



Continuous impact on future Windows versions





Conclusion

- The first formulation and real-world deployment of performance debugging in the large
 - as a data mining problem on callstacks
- A mining-clustering mechanism for reducing costly-pattern mining results
 - based on domain-specific characteristics of callstacks
- Industrial impact on using StackMine in performance debugging in the large for Microsoft Windows





Acknowledgment

- Our partners in Microsoft product teams
- The researchers from Microsoft Research





Q&A

Thank you!